



OST

Eastern Switzerland
University of Applied Sciences

Efficient Public Transit Routing with RAPTOR

A Production-Ready Solution Leveraging GTFS Data

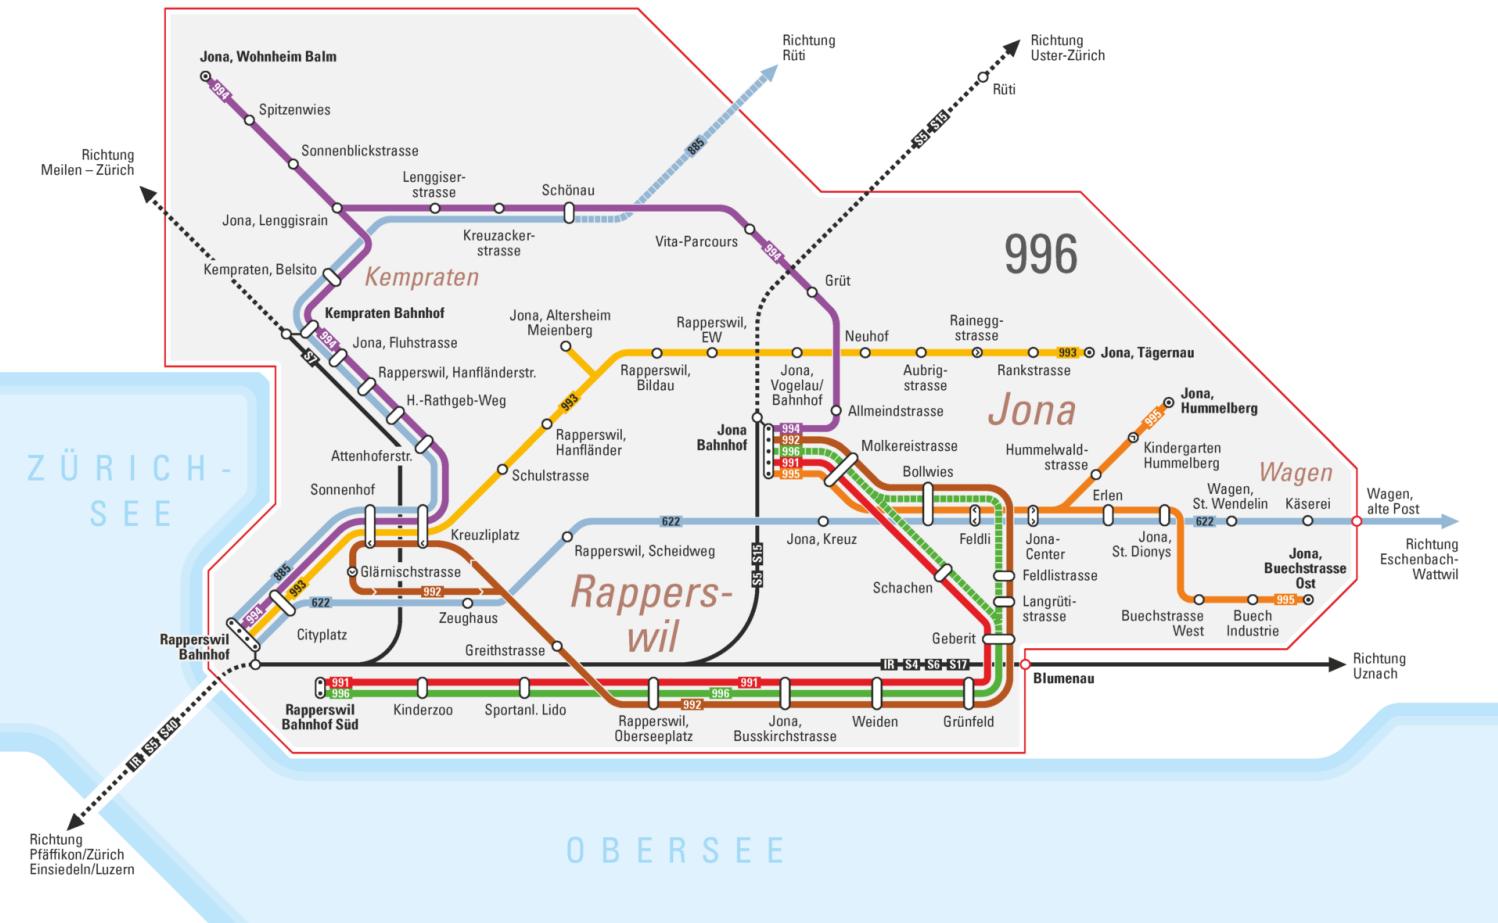
Michael Brunner, Lukas Connolly, Merlin Unterfinger
Supervisor: Thomas Letsch

24 October 2024

MAS Software Engineering 2022 - 2024

Agenda

- Project Management
 - Analysis
 - Design
 - Implementation
 - Demo
 - Results
 - Conclusion



Source: <https://www.vzo.ch/stadtbus-rapperswil-jona>

Project Management

Team



Lukas Connolly
Senior Research Engineer

Zimmer Biomet
Computational Biomechanics
Sulzerallee 8
CH-8404 Winterthur

lukas.connolly@zimmerbiomet.com



Michael Brunner
Software Engineer

Cadwork Holz AG
Informatik
Industriestrasse 28
CH-9100 Herisau

brunner@cadwork.swiss



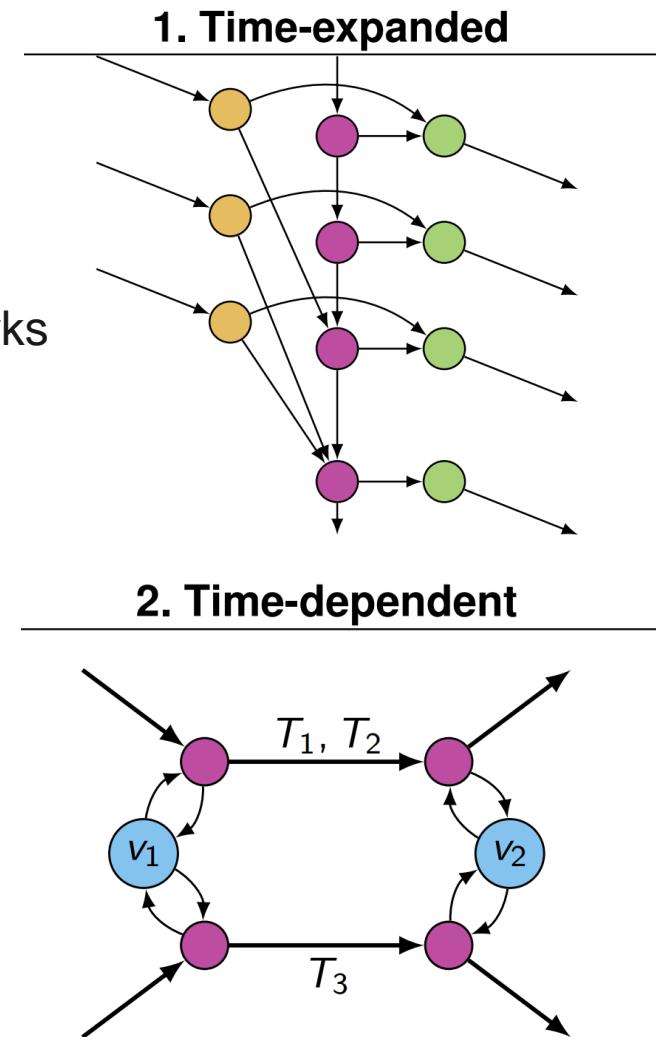
Merlin Unterfinger
Subject Matter Expert

SBB AG
Corporate Development
Hilfikerstrasse 1
CH-3000 Bern 65

merlin.unterfinger@sbb.ch

Motivation

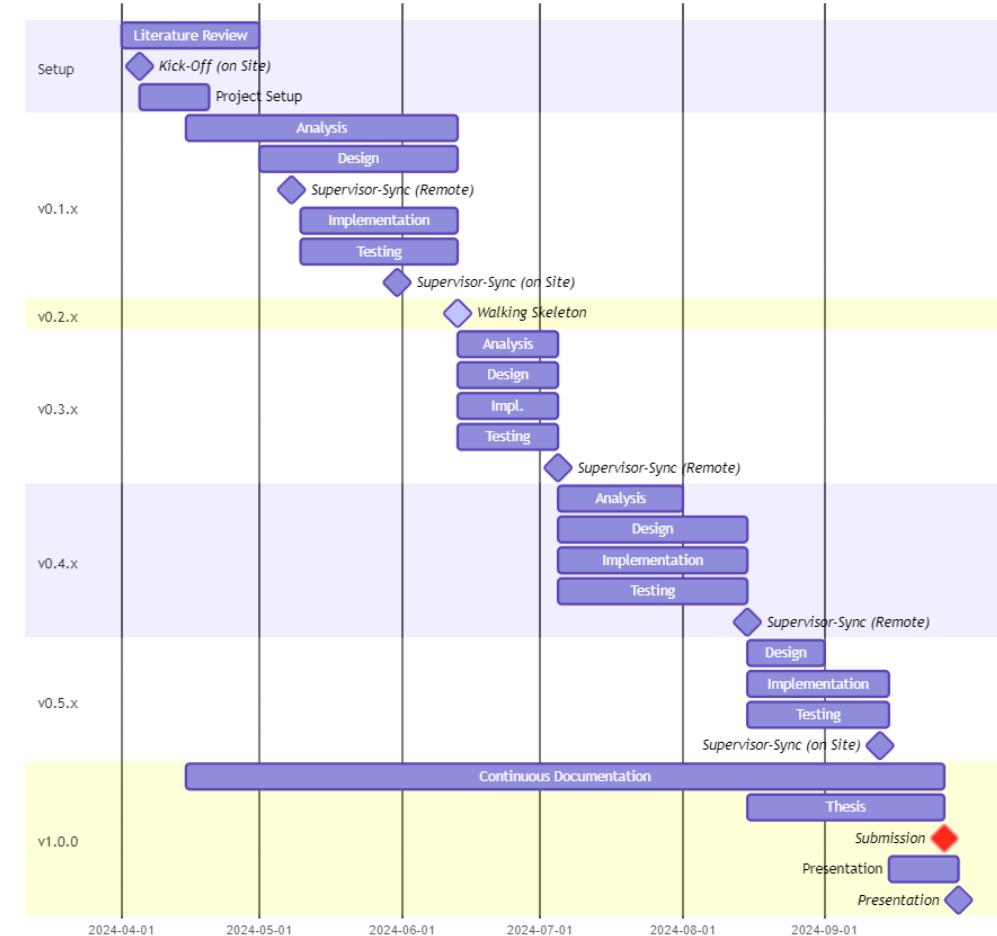
- Public transit routing is a **challenging topic**
 - Time-dependent transit schedules add complexity
 - Traditional graph-based algorithms struggle with large-scale networks
 - The RAPTOR algorithm offers efficient, low-preprocessing routing
- **Open Data** availability
 - General Transit Feed Specification (GTFS) standard
 - Many national transit agencies publish timetables
- Project integrates **software engineering** skills
 - Algorithms and data structures, architecture
 - Language-specific optimizations (Java, C++)



Source: Jonas Sauer (2021): Public Transit Journey Planning

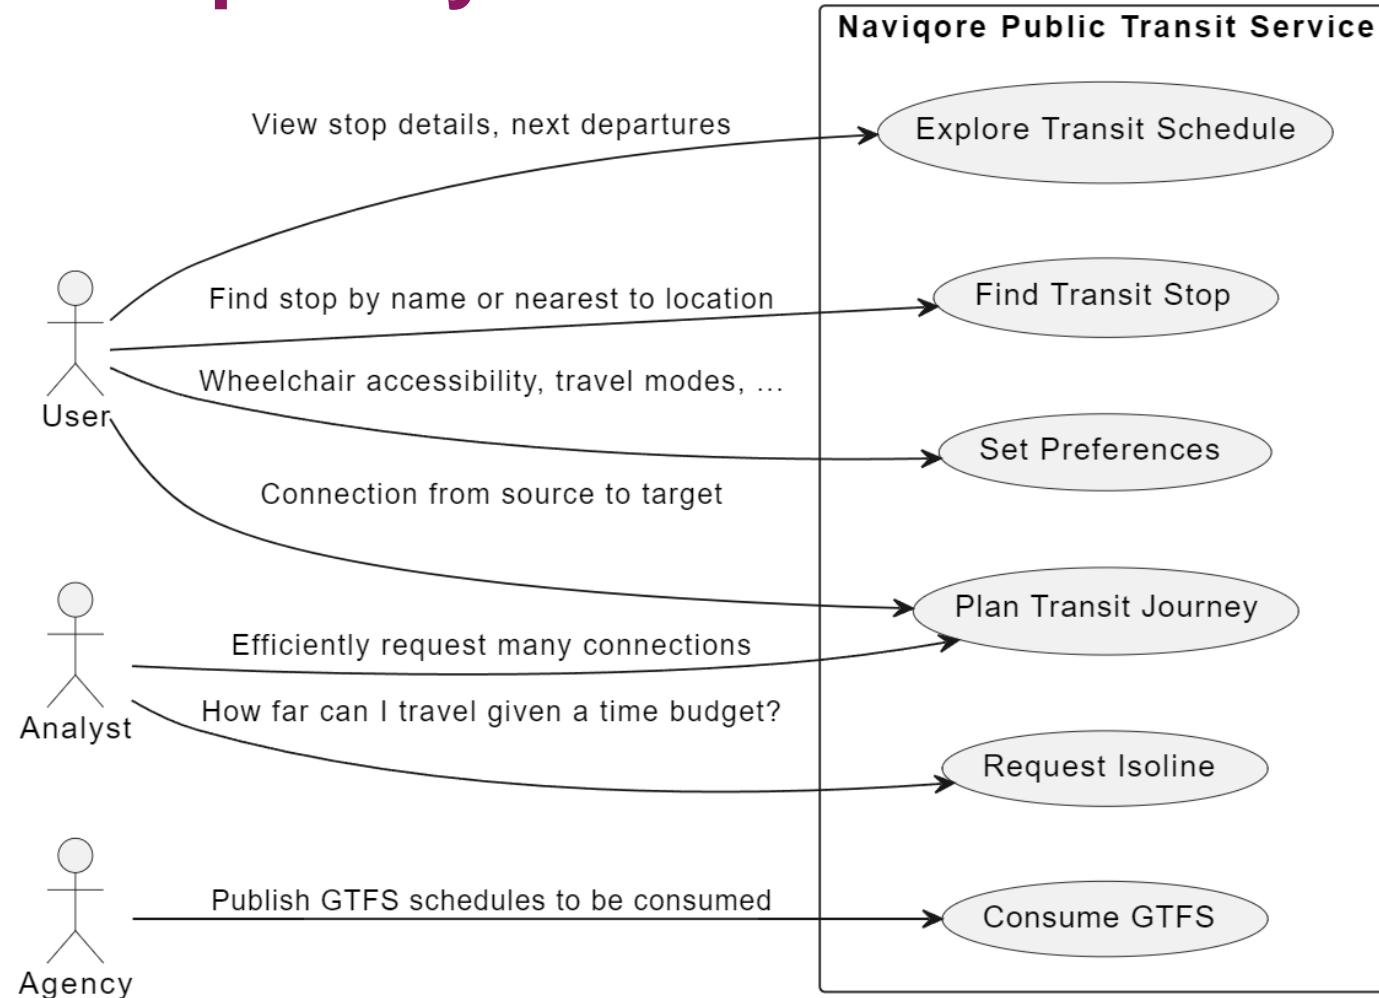
Software Development Process

- Project structured in **3 phases**
 - Walking skeleton (*v0.2.0*)
 - Main implementation (*v1.0.0*)
 - Feature freeze
- **Scrum-like** development setup
 - Sprints of 1 month
 - Weeklies, planning, demo & review
 - Task and time management in Jira
- **Written guidelines**
 - Code style, Git & release process
 - Comprehensive pull request reviews





Use Cases – Naviqore System





Requirements – Naviqore System

- Structure: {TYPE}-{COMPONENT}-{PRIORITY}{NUMBER}

Requirement Types

UC: Use case

NF: Non-function requirement

Components

SE: Service

SC: Schedule

RO: Routing

VW: Viewer

Priority Levels

M: Must-have

S: Should-have

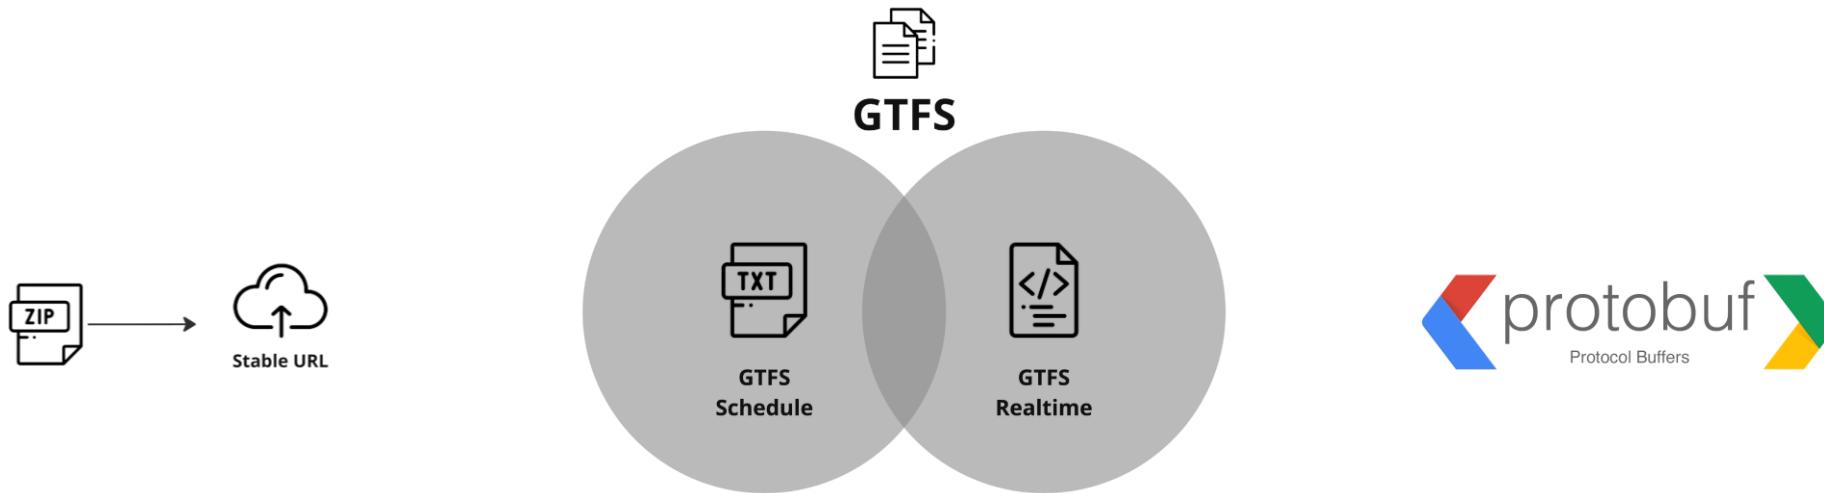
N: Nice-to-have

- All must-have, and most of should-have requirements were satisfied

	Service	Schedule	Routing	Viewer	Total
Must-have	7 (7)	4 (4)	5 (5)	1 (1)	17 (17)
Should-have	4 (4)	1 (0)	7 (6)	1 (1)	13 (11)
Nice-to-have	4 (1)	1 (0)	4 (1)	7 (2)	16 (4)
Total	15 (12)	6 (4)	16 (12)	9 (4)	46 (32)

Analysis

GTFS (General Transit Feed Specification)



- Text files
- **Fixed** schedules, stops, and routes
- Timetables, stop locations, fare information
- Protocol buffer
- **Live** updates on transit services
- Delays, vehicle locations, service alerts.

Source: <https://opentransportdata.swiss/en/cookbook/gtfs/>

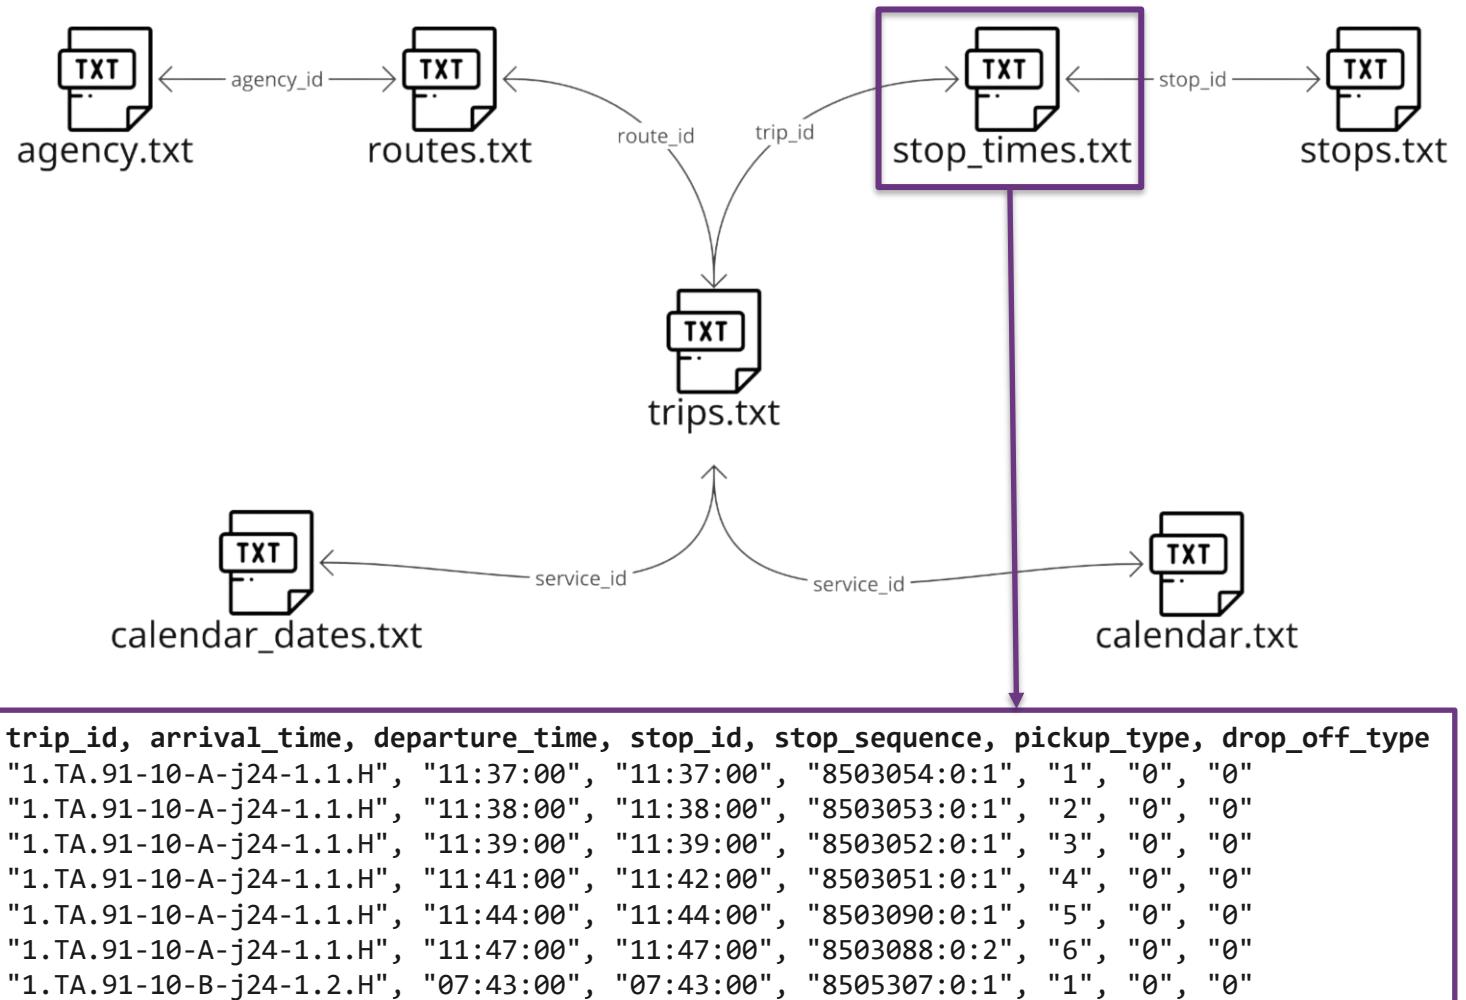
Analysis

GTFS – Static

Timetable of Switzerland:

File	Size
agency.txt	43.0 KB
calendar.txt	3.3 MB
calendar_dates.txt	194.0 MB
routes.txt	208.0 KB
stop_times.txt	1.4 GB
stops.txt	6.4 MB
transfers.txt	1.5 MB
trips.txt	136.0 MB
Total	1.7 GB

Source: <https://opentransportdata.swiss/en/group/timetables-gtfs>



Source: https://gtfs.org/getting_started/create/

Analysis

RAPTOR Data Structure

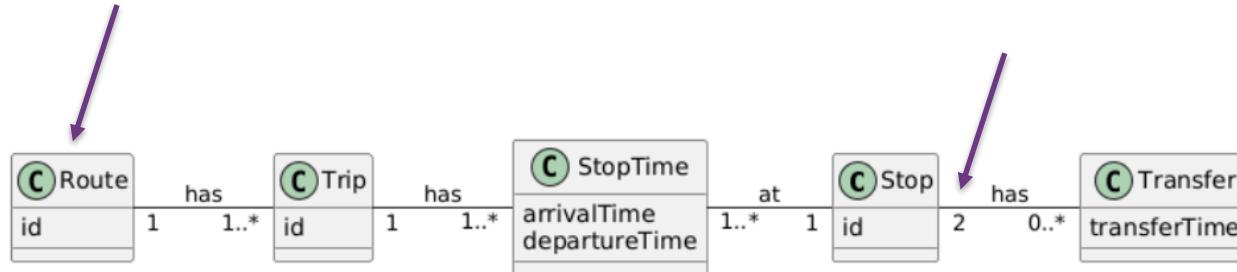


Figure 3.4: Domain model for RAPTOR.

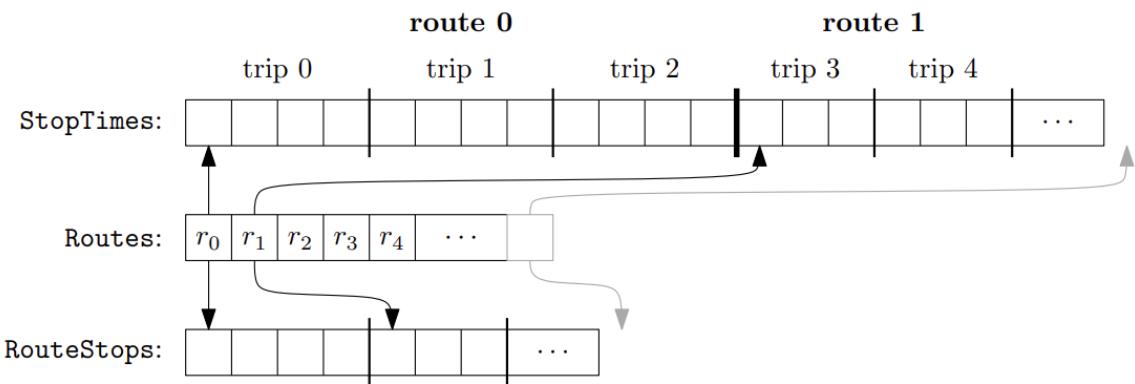


Figure 3: Illustration of the adjacency structure of routes.

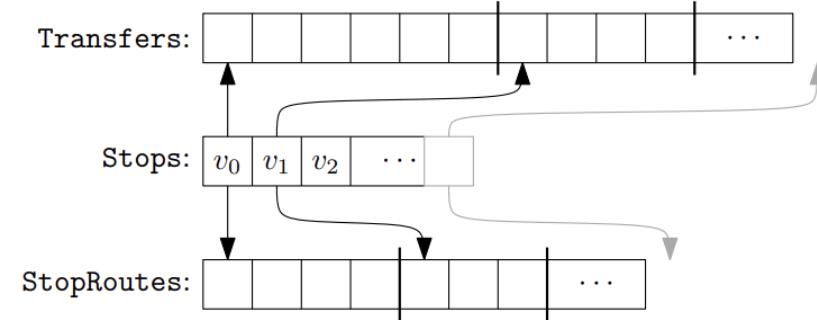


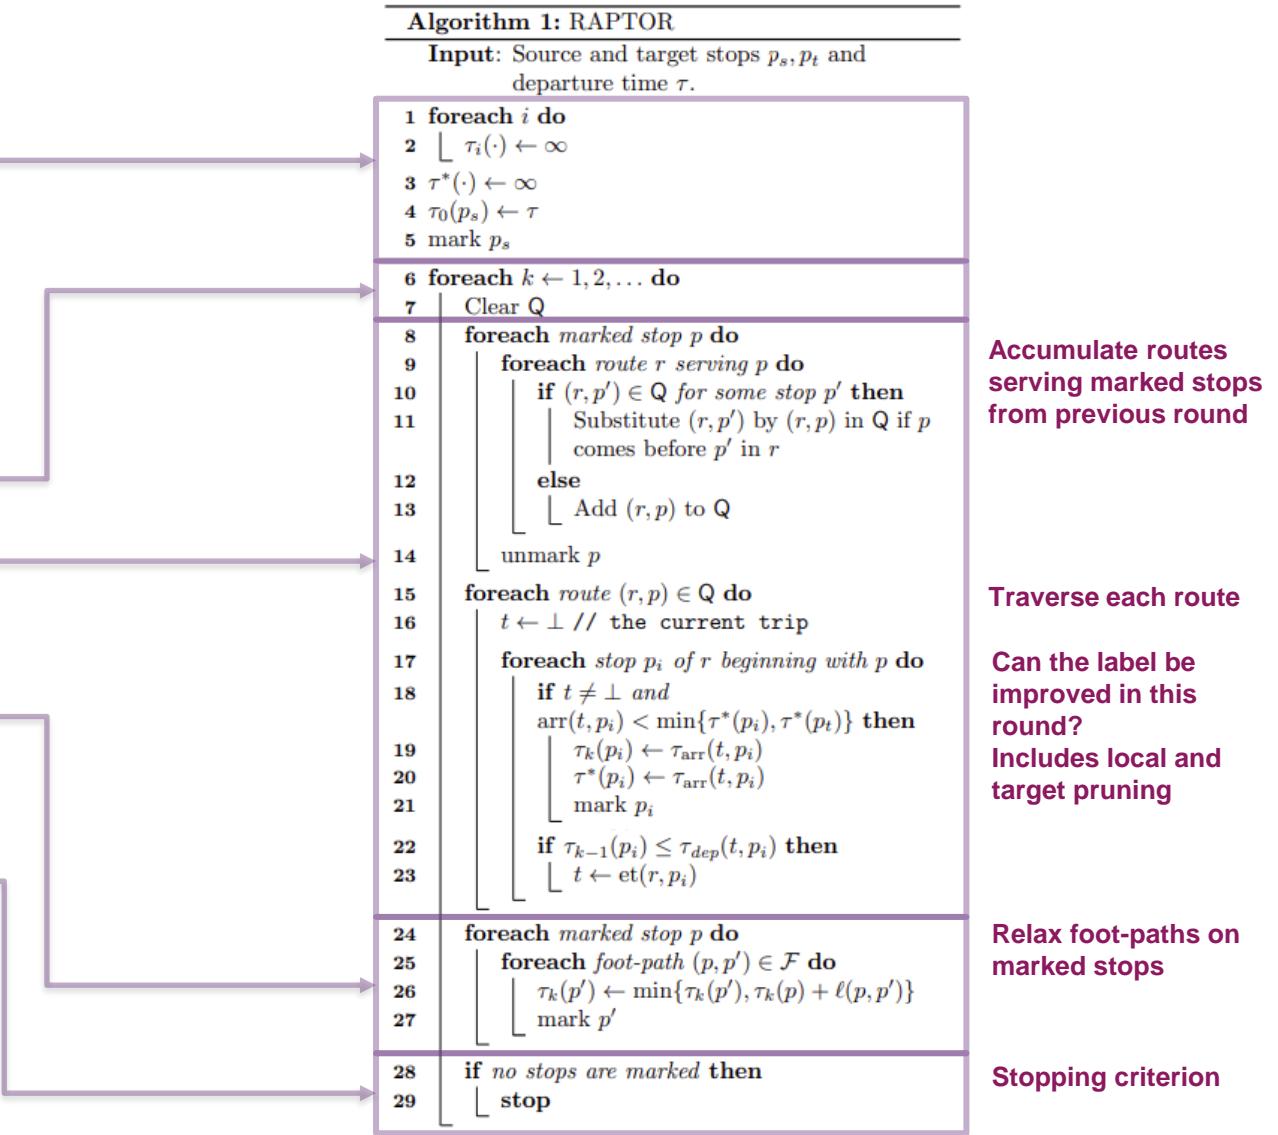
Figure 4: Illustration of the adjacency structure of stops.

Source: <https://doi.org/10.1137/1.9781611972924.13>

Analysis

RAPTOR Algorithm

- **Initialization:** Mark first stop
- **Initial footpath relaxation**
- **Main Loop: Round-based search**
 - Add new round
 - **Scan** routes, mark improved stops
 - **Relax** footpaths, mark improved stops
- **Termination:** End when no improvements possible
- **Reconstruct:** Best route

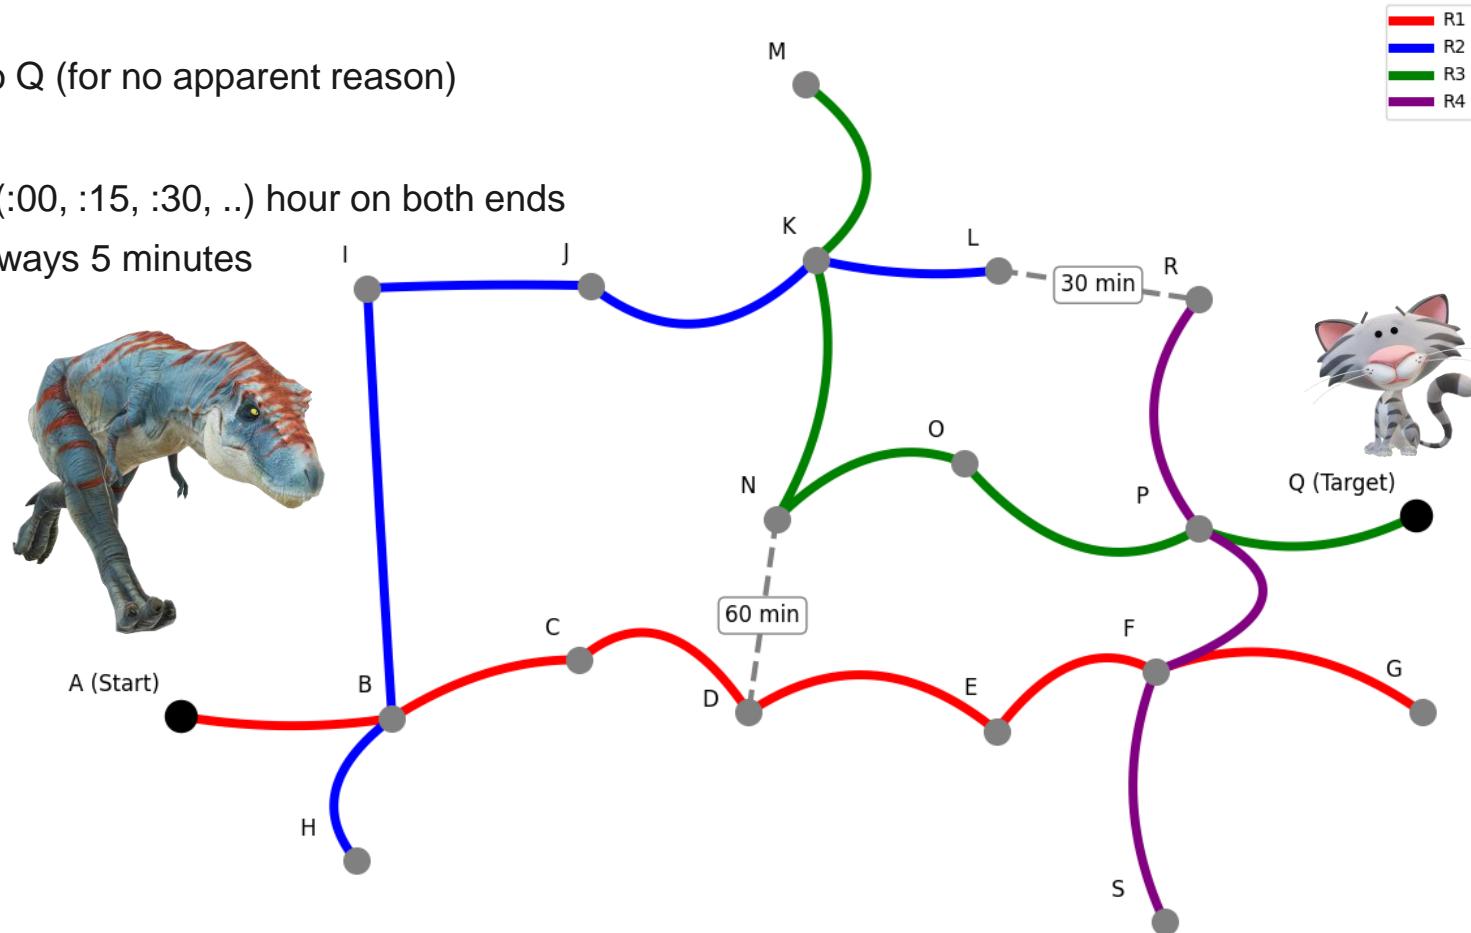


Source: <https://doi.org/10.1137/1.9781611972924.13>

Analysis

RAPTOR Algorithm

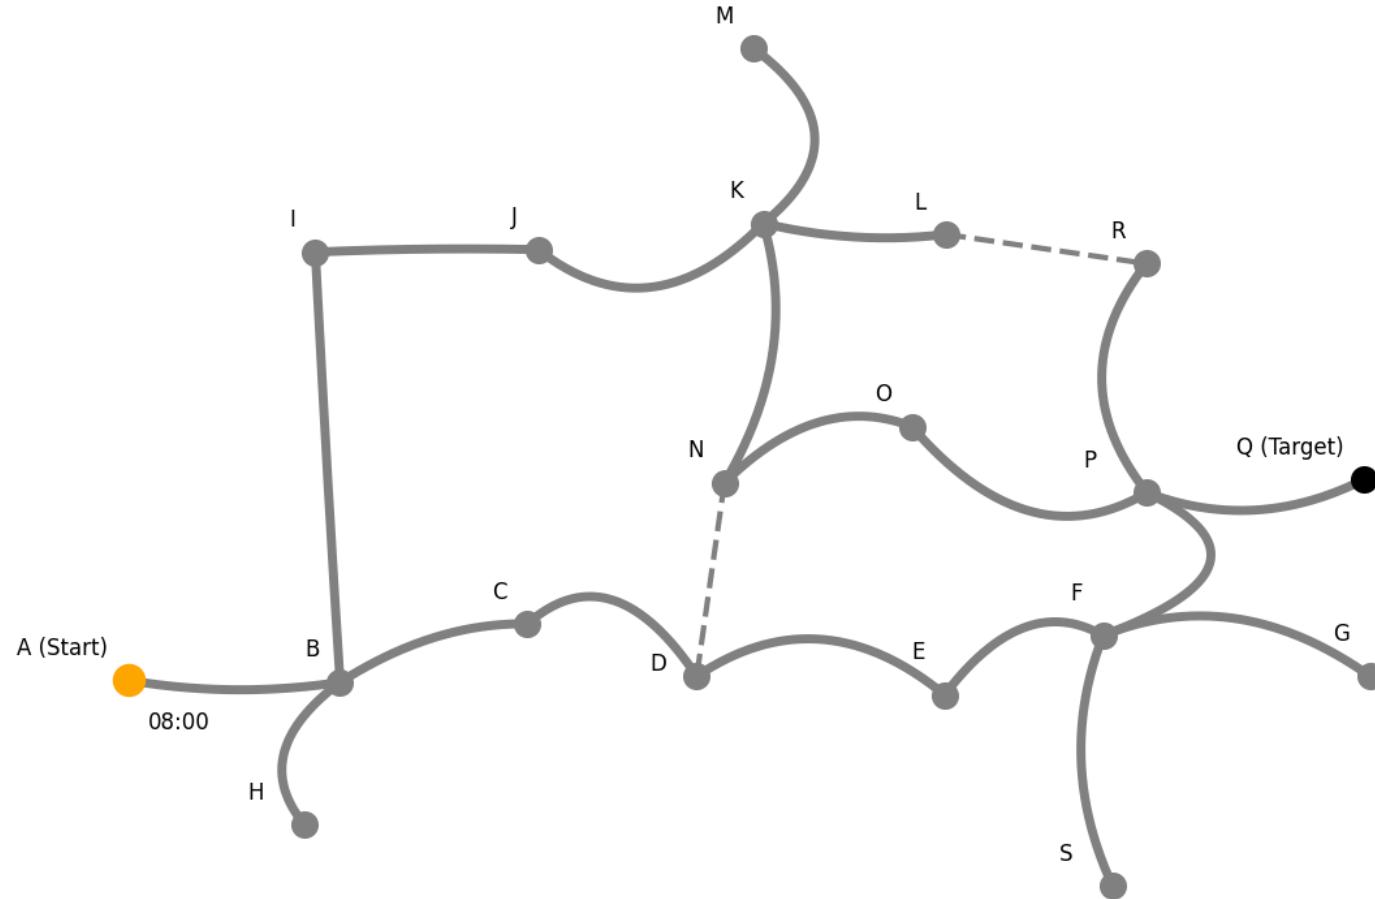
- The Raptor wants to get from A to Q (for no apparent reason)
- The current time is 08:00 AM
- All routes start every 15 minutes (:00, :15, :30, ..) hour on both ends
- The time between each stop is always 5 minutes



Analysis

RAPTOR Algorithm (Round 0)

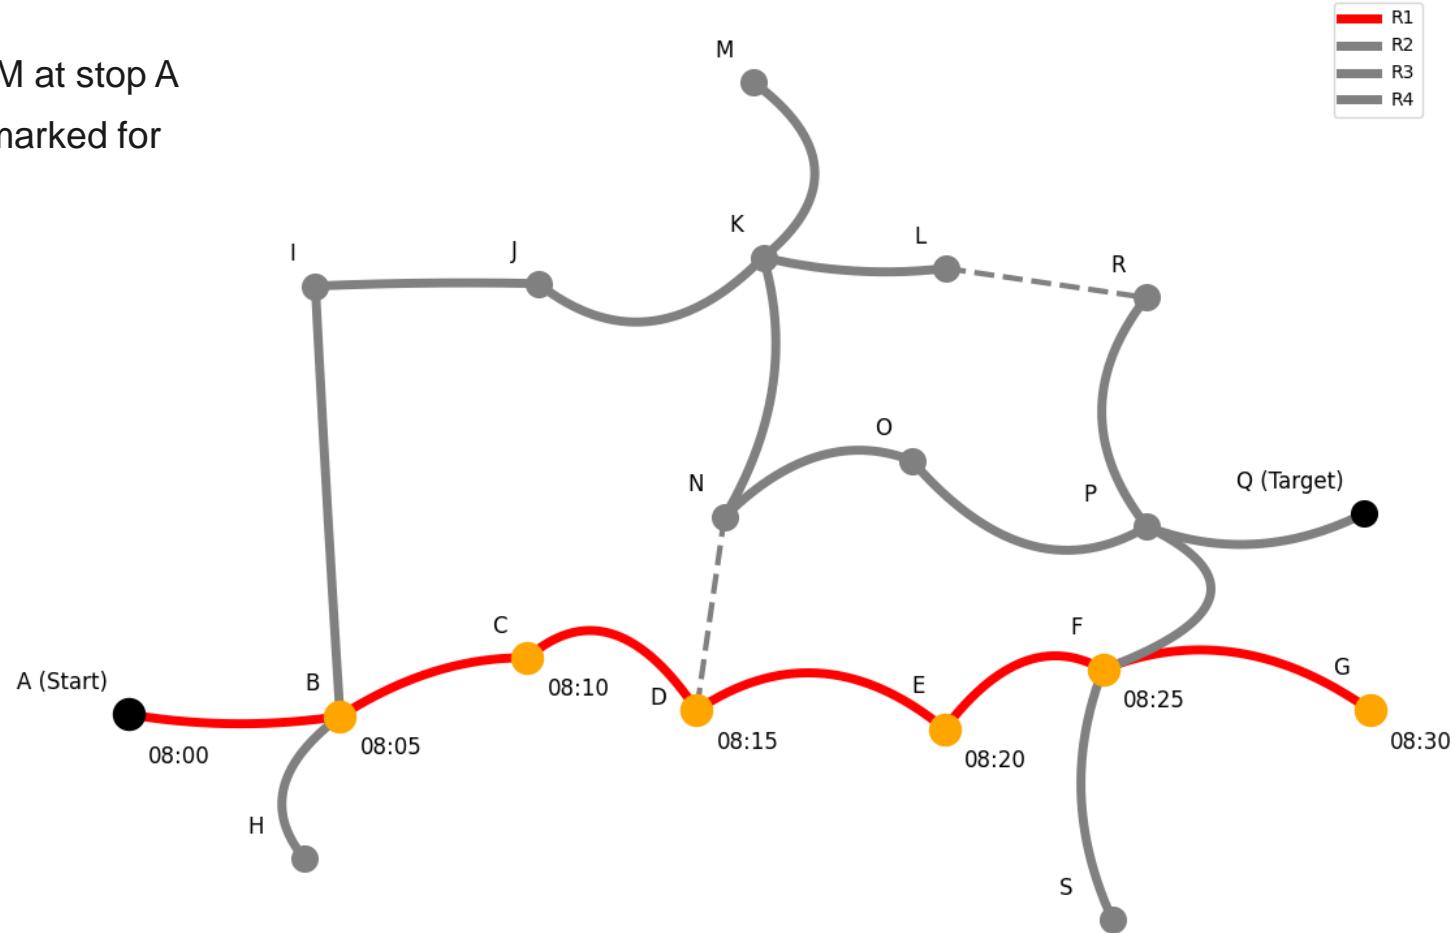
- Stop A is marked
- No transfers are possible from A



Analysis

RAPTOR Algorithm (Round 1 – Route Scanning)

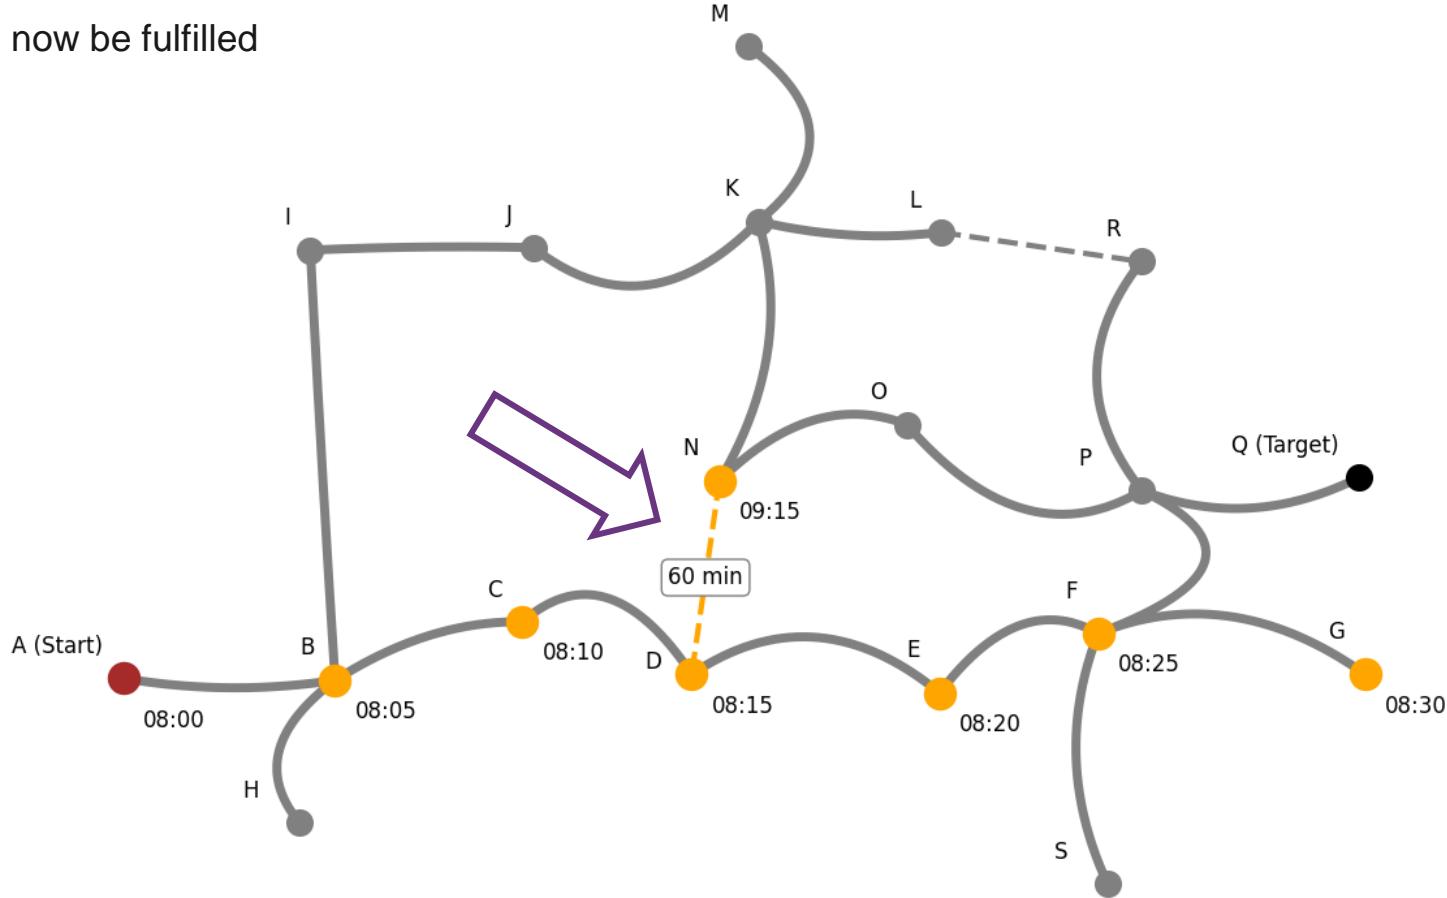
- Route R1 can be entered at 08:00 AM at stop A
- All stops along the line can now be marked for next round scanning



Analysis

RAPTOR Algorithm (Round 1 – Footpath Relaxation)

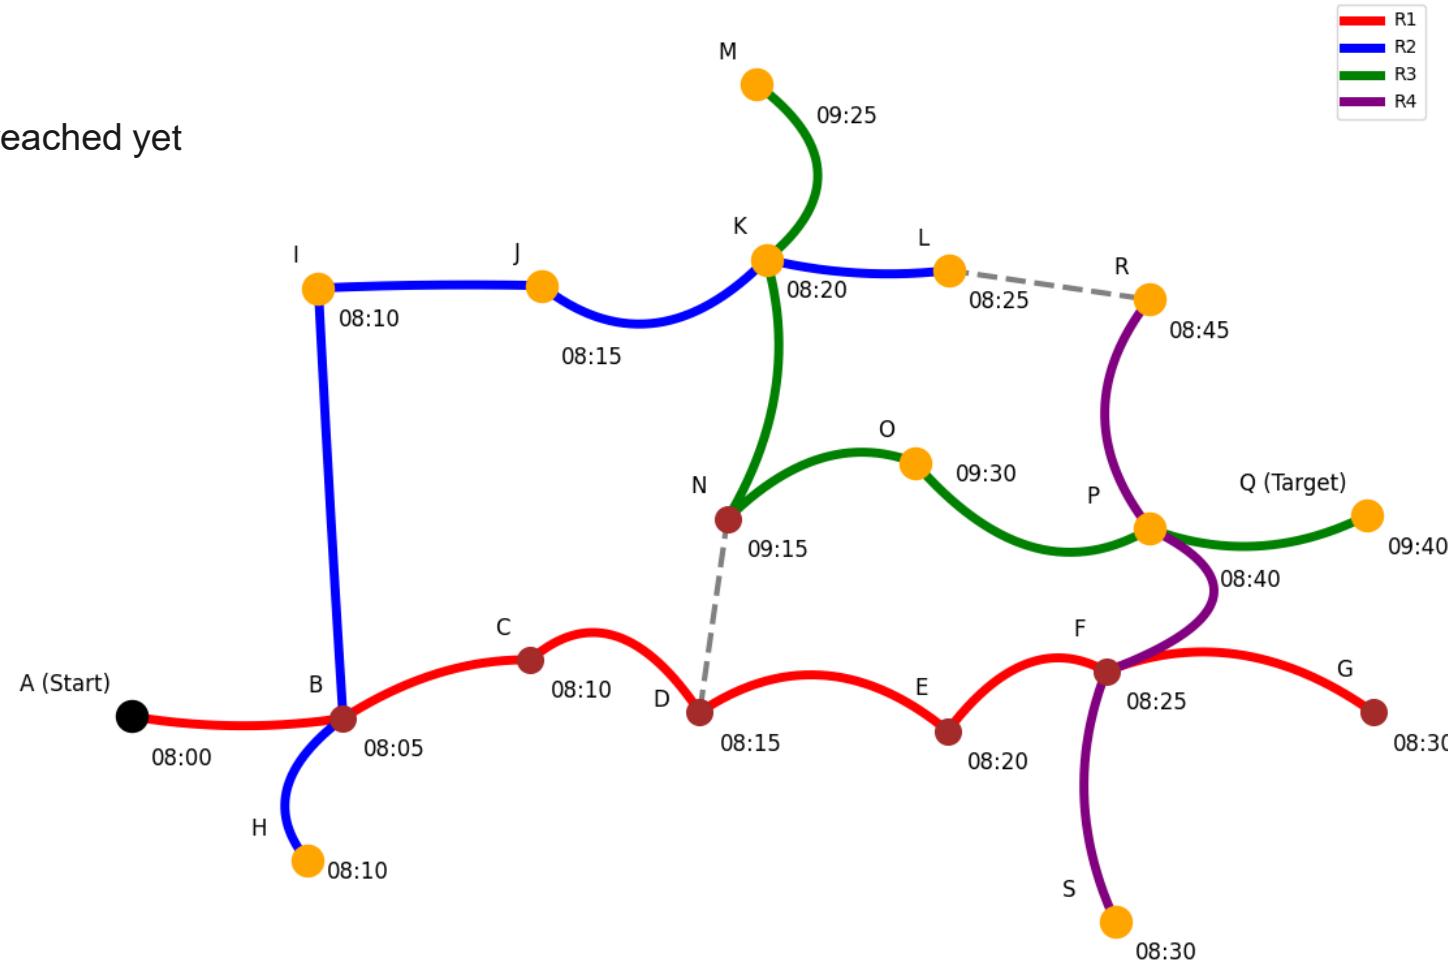
- Stop D has a transfer to N which can now be fulfilled



Analysis

RAPTOR Algorithm (Round 2 – Route Scanning)

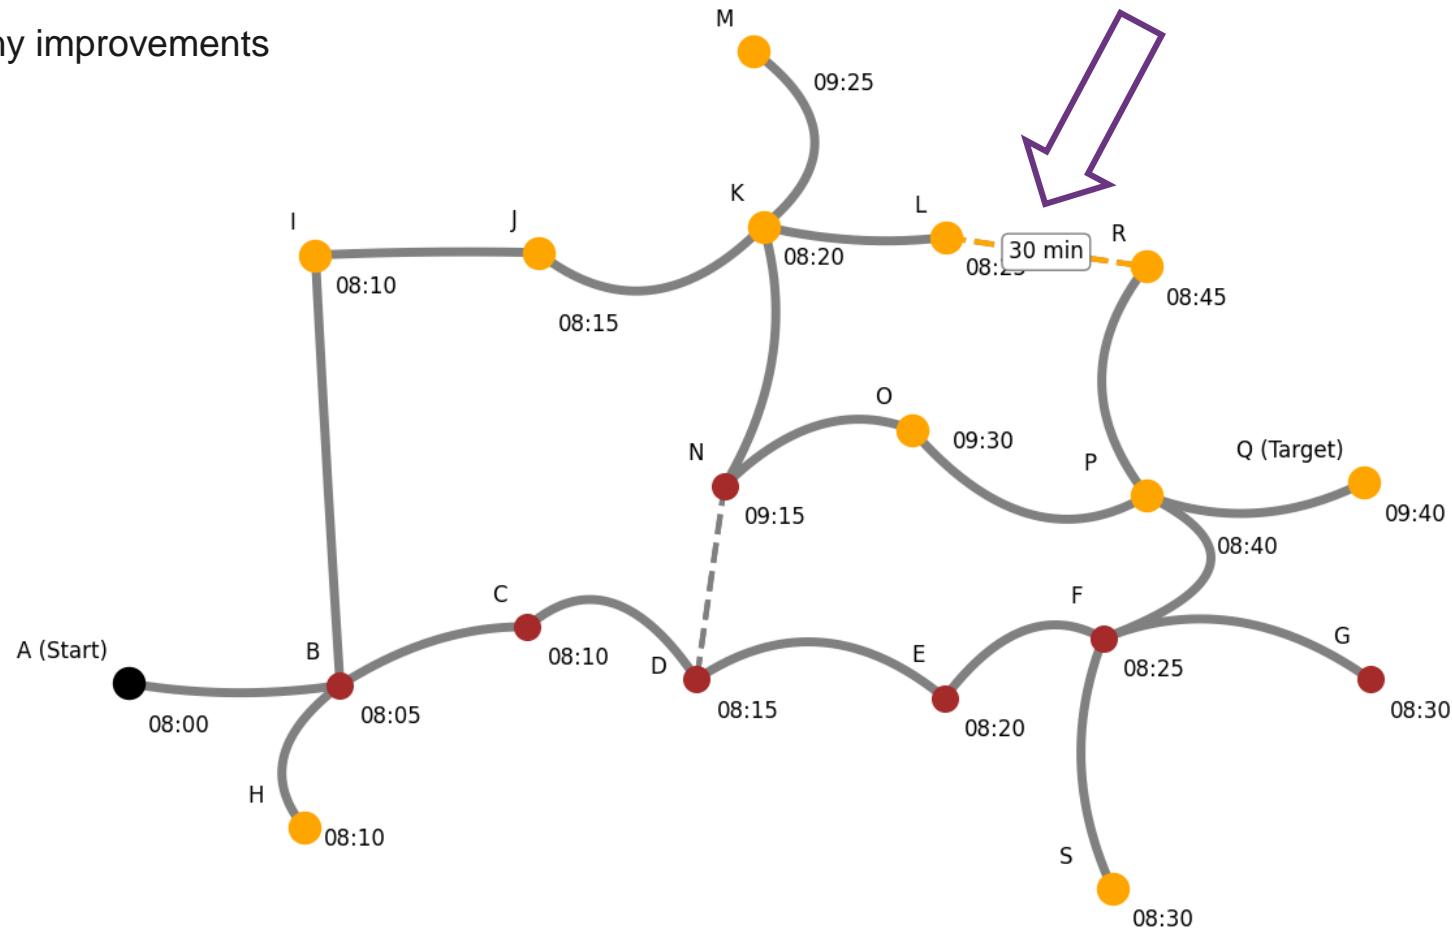
- Now all routes need to be scanned
- Marking all stops that haven't been reached yet



Analysis

RAPTOR Algorithm (Round 2 – Footpath Relaxation)

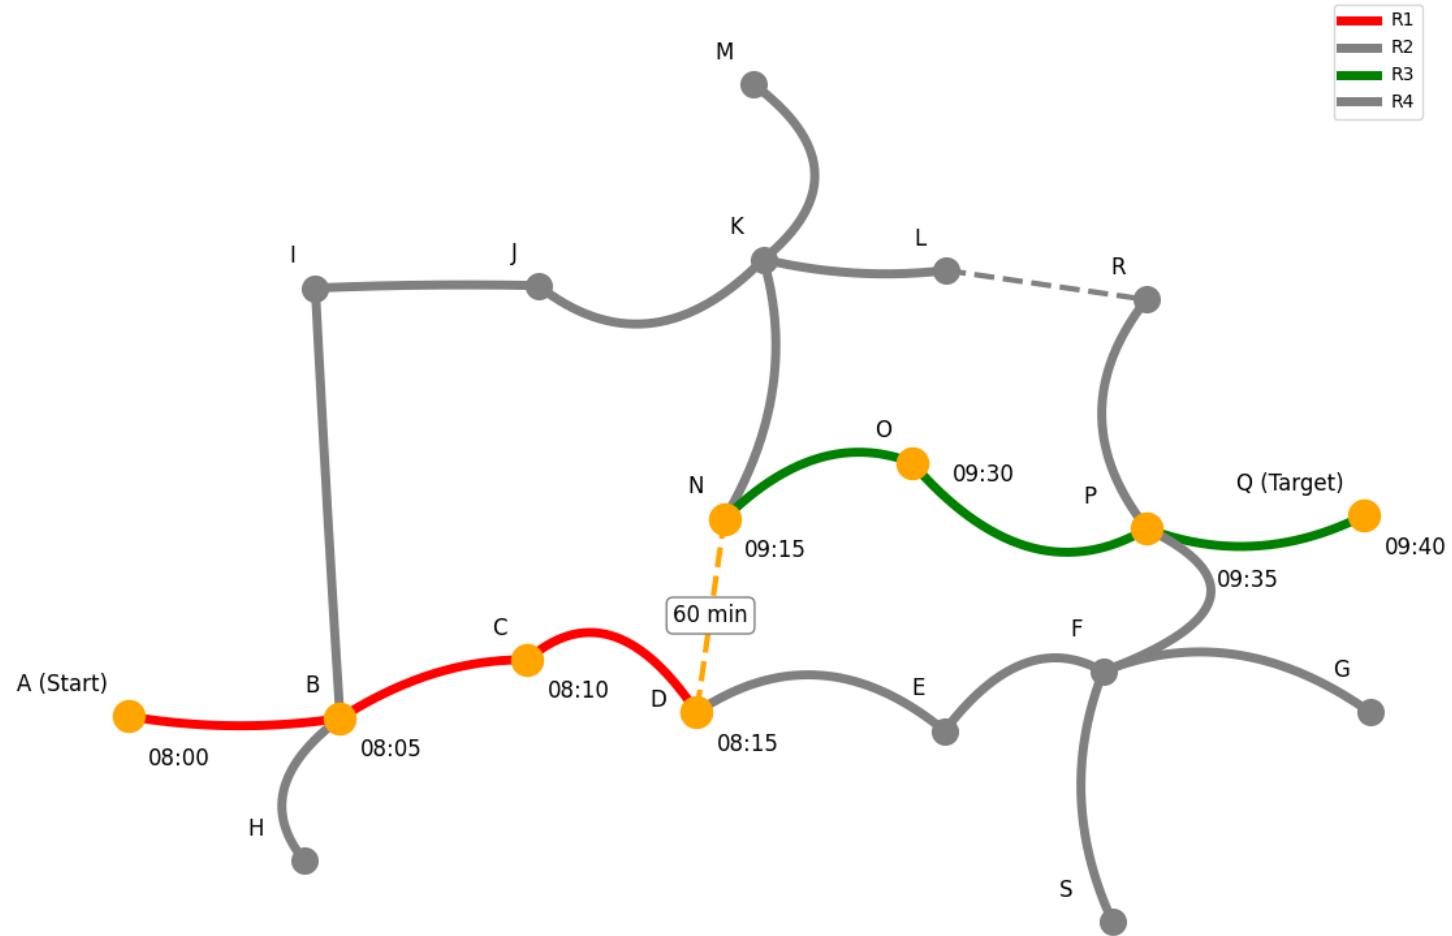
- Transfer from L to R does not add any improvements



Analysis

RAPTOR Algorithm

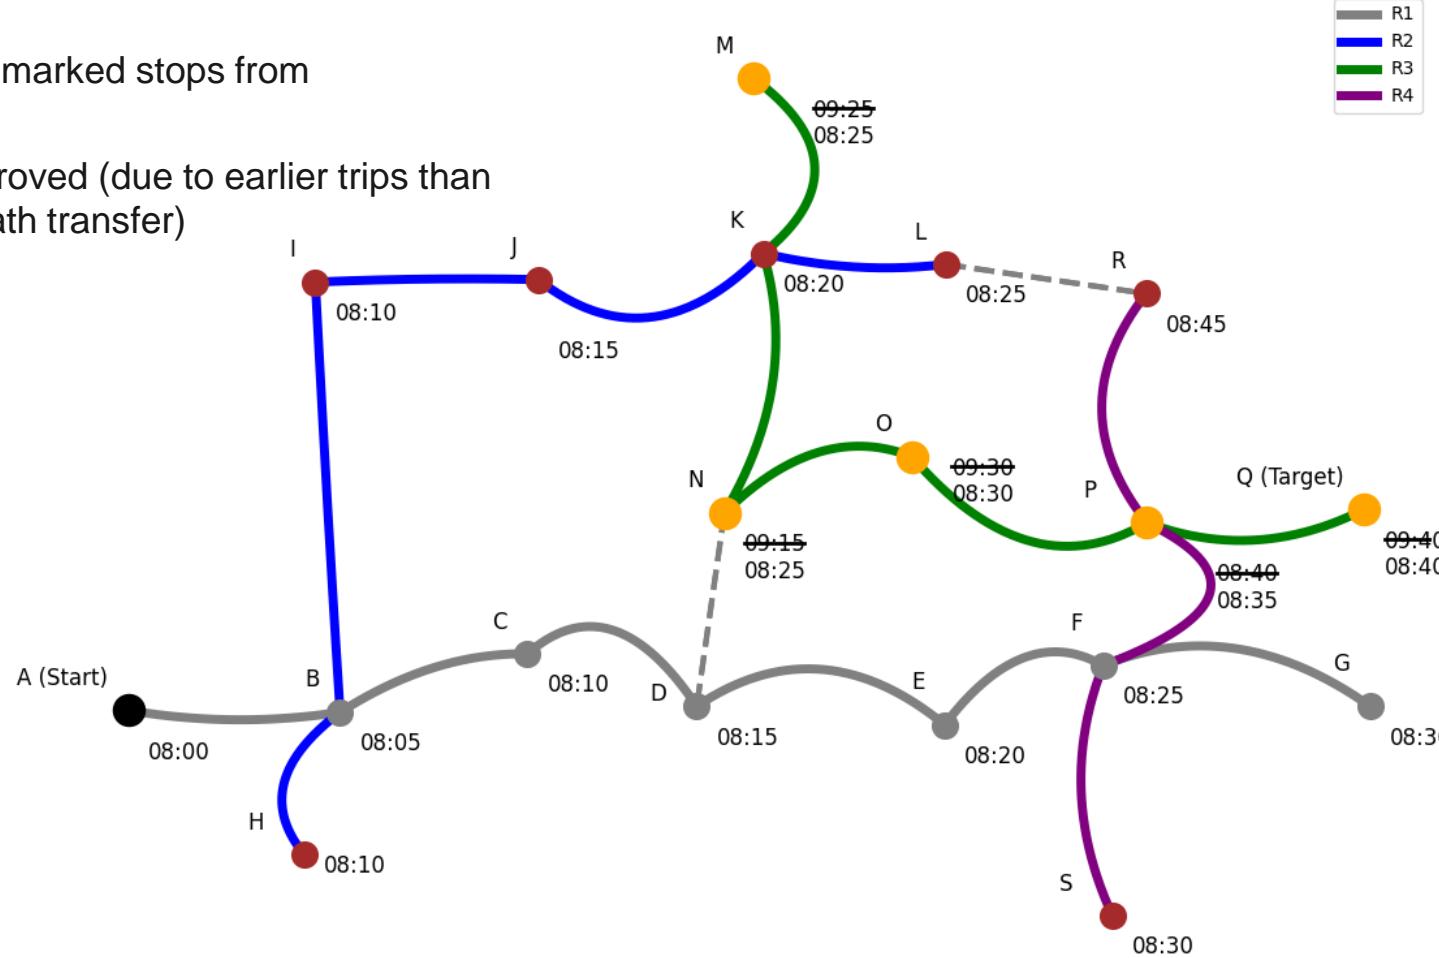
- First Pareto-optimal connection



Analysis

RAPTOR Algorithm (Round 3 – Route Scanning)

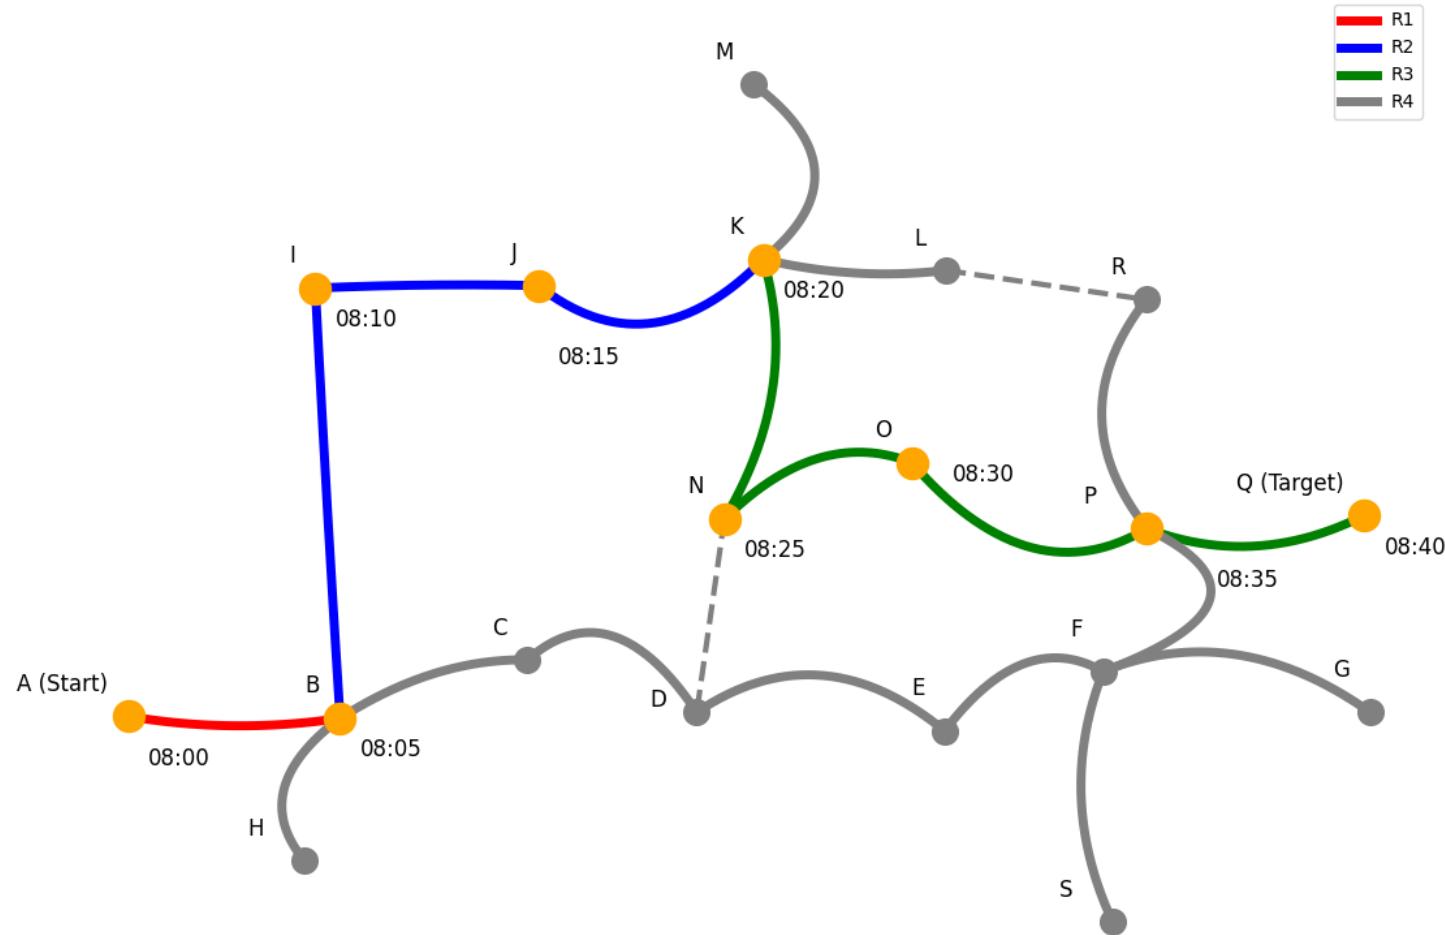
- Scanning all routes passing through marked stops from last round (R2, R3 and R4)
 - Only stops on R3 can be further improved (due to earlier trips than previously achievable through footpath transfer)

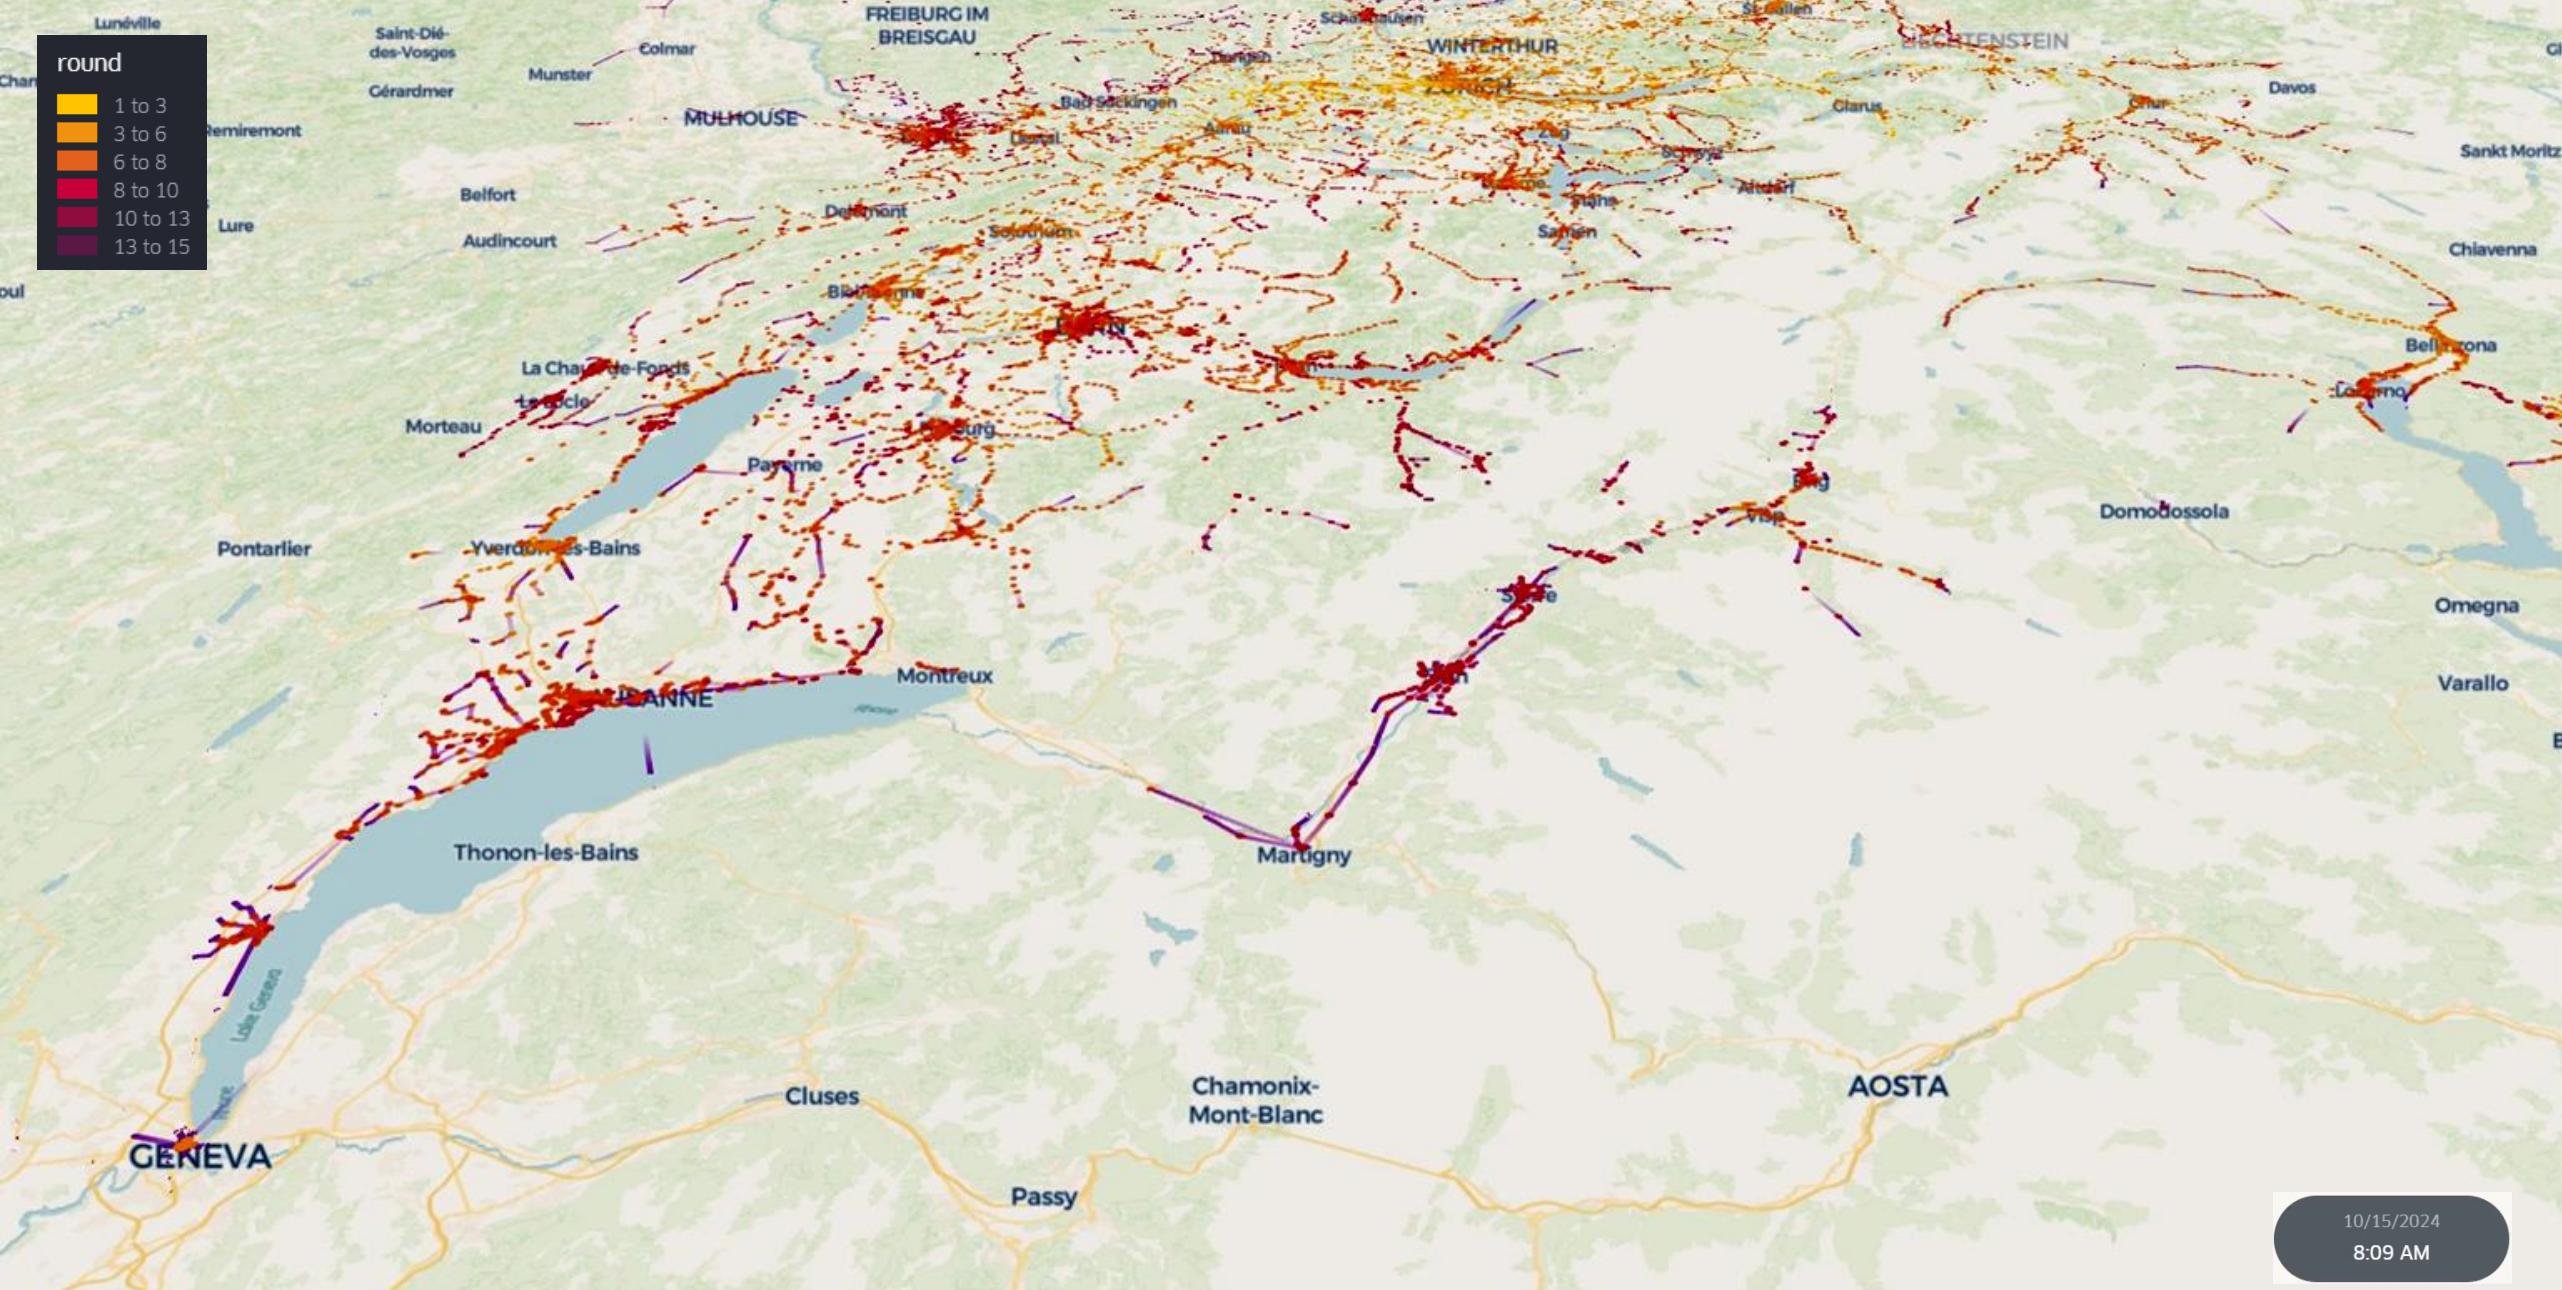


Analysis

RAPTOR Algorithm

- Second Pareto-optimal connection:
More rounds but earlier arrival time

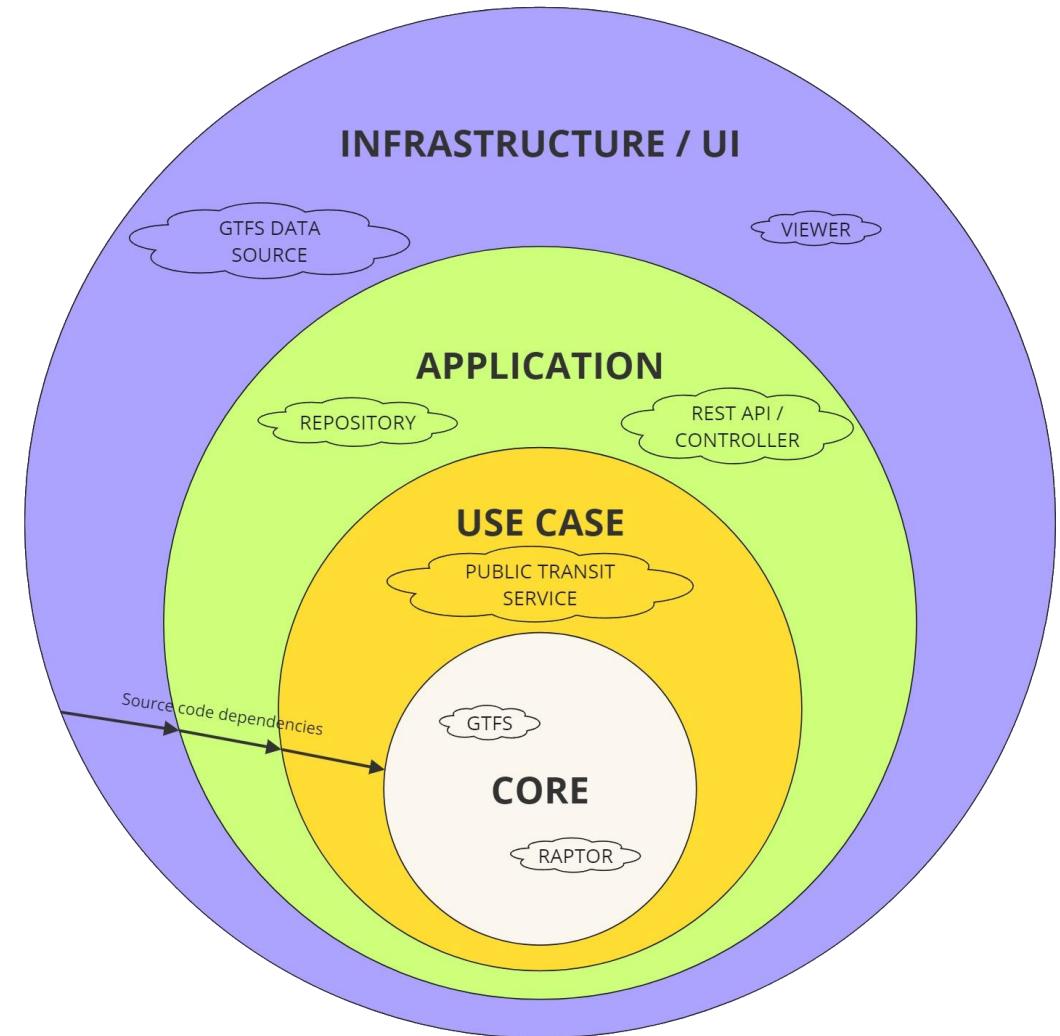




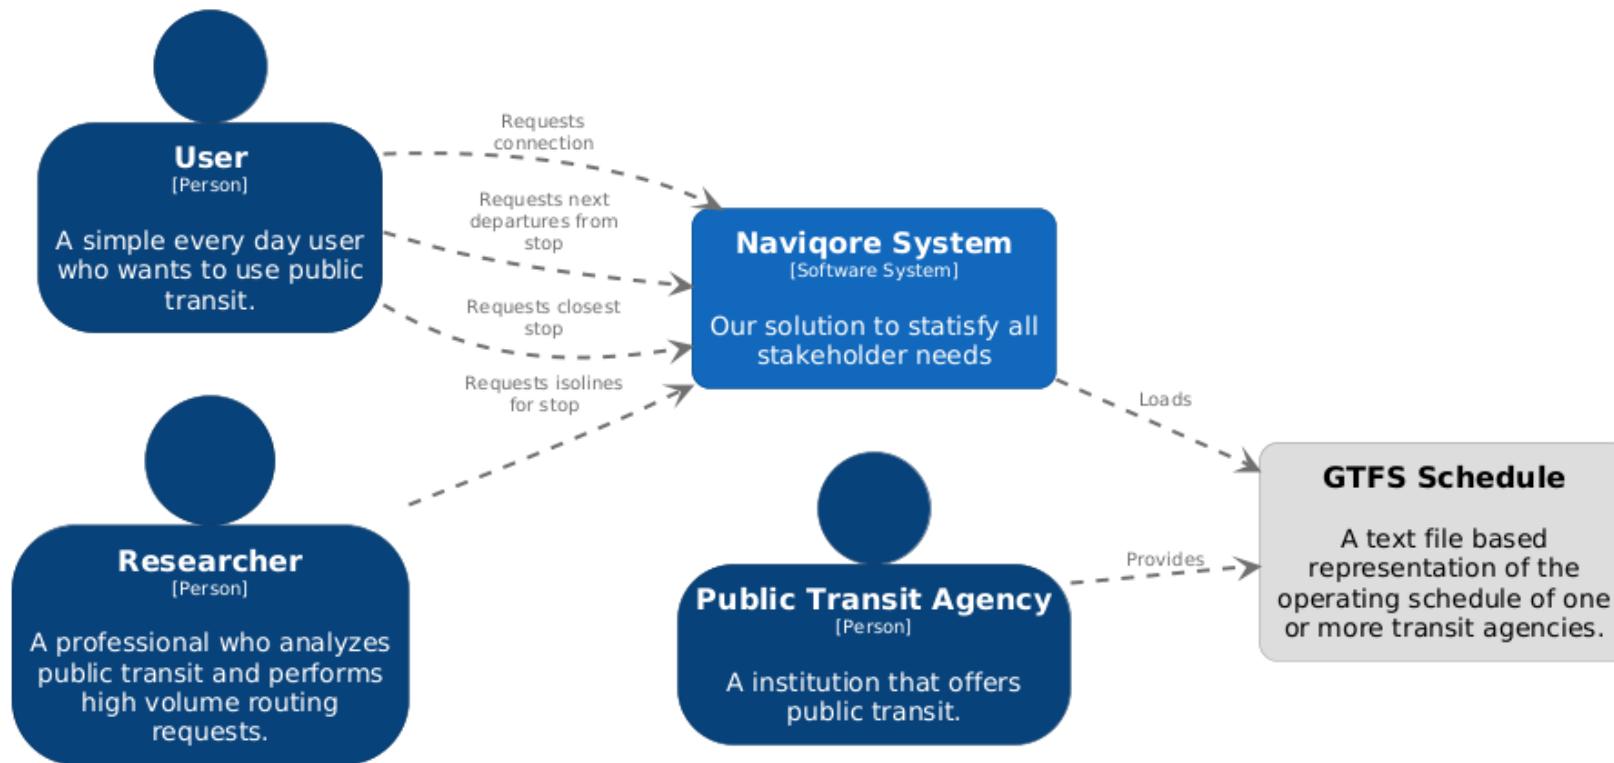
10/15/2024
8:09 AM

Architecture

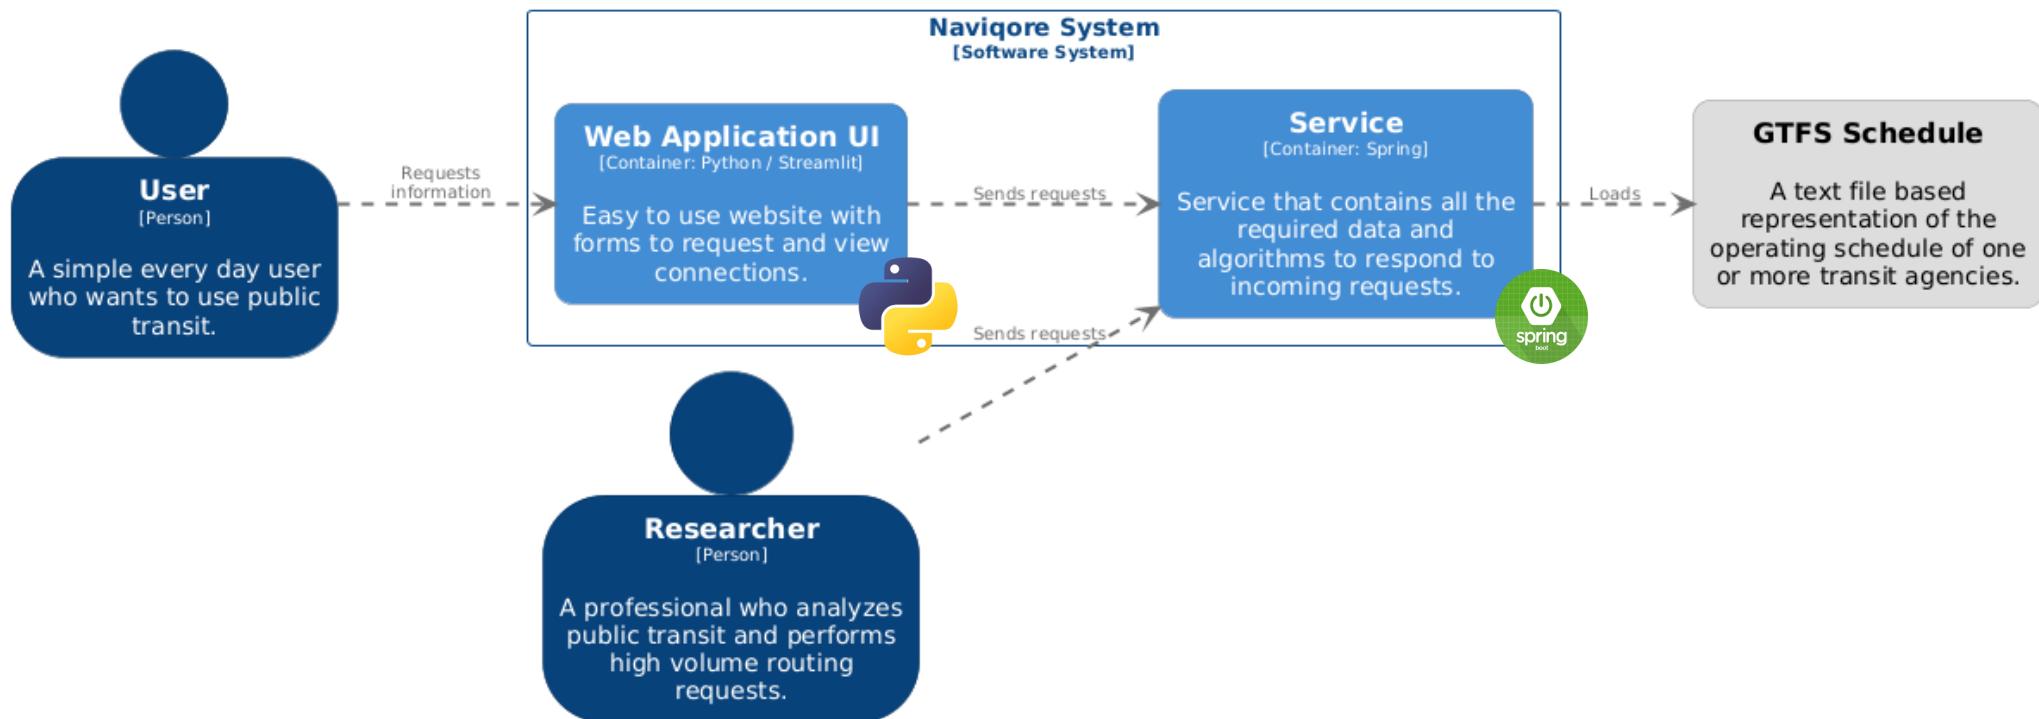
- Clean and **Onion** Architecture.
- **Four layers**
 - Core, Use Case, App, Infrastructure/UI
- Inward dependencies via **dependency injection**
- Inner **abstractions**, outer implementations
- **Loose coupling**, replaceable implementations
- App layer isolates
 - Spring Boot components
 - Repository: Timetable data source



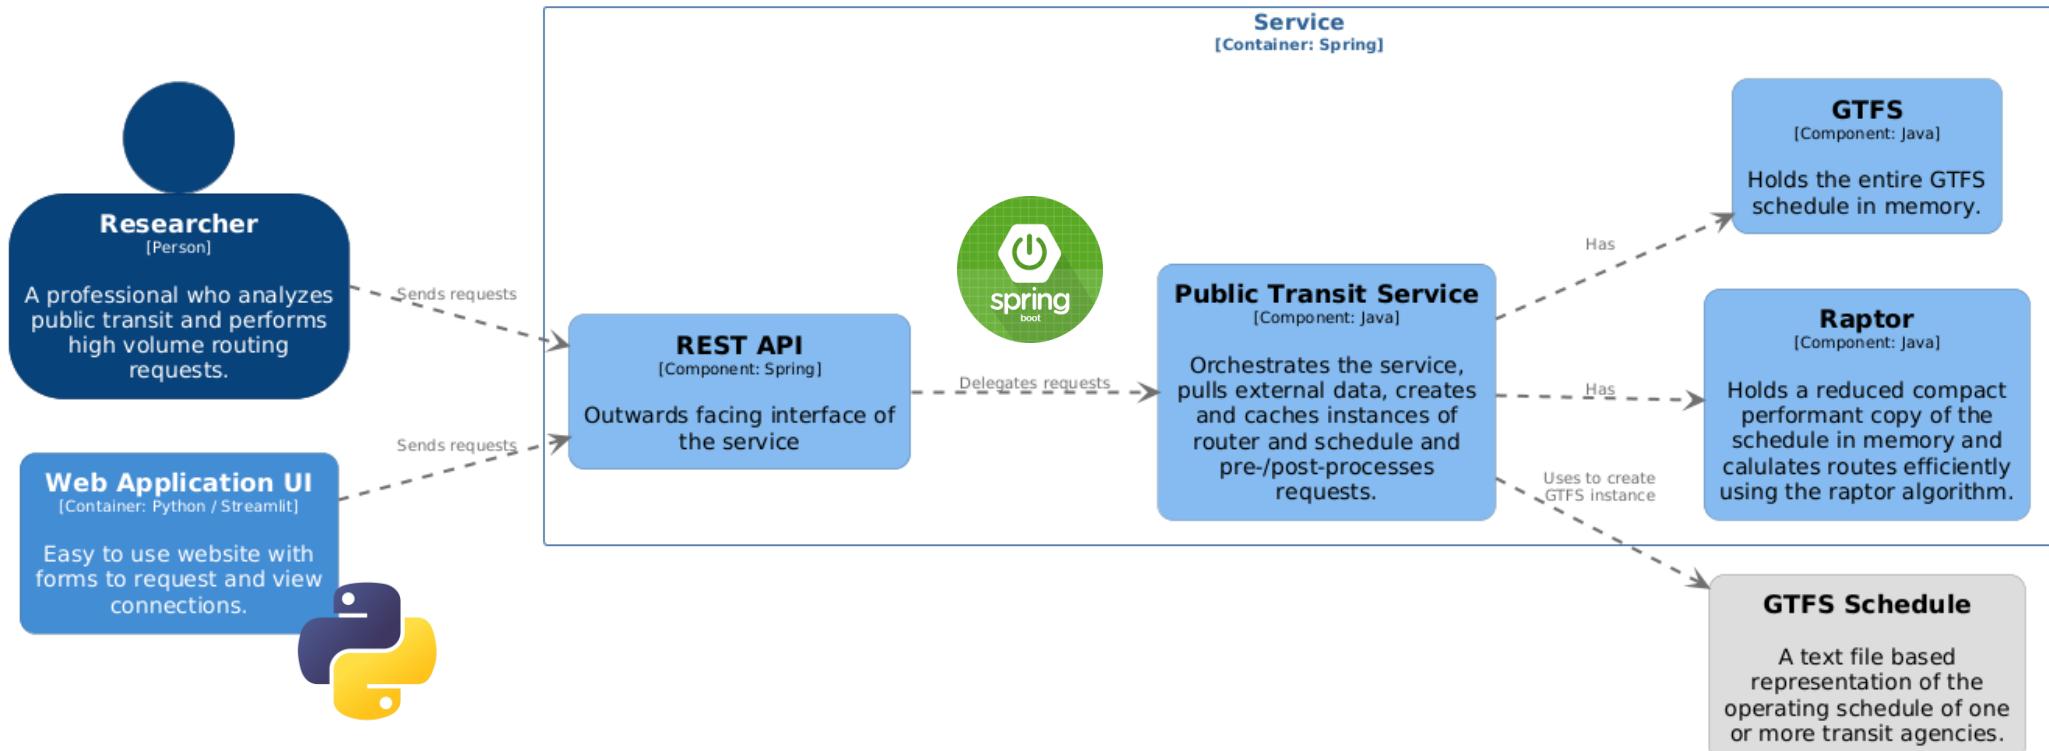
C4 Model – System Context



C4 Model - Containers



C4 Model – Components



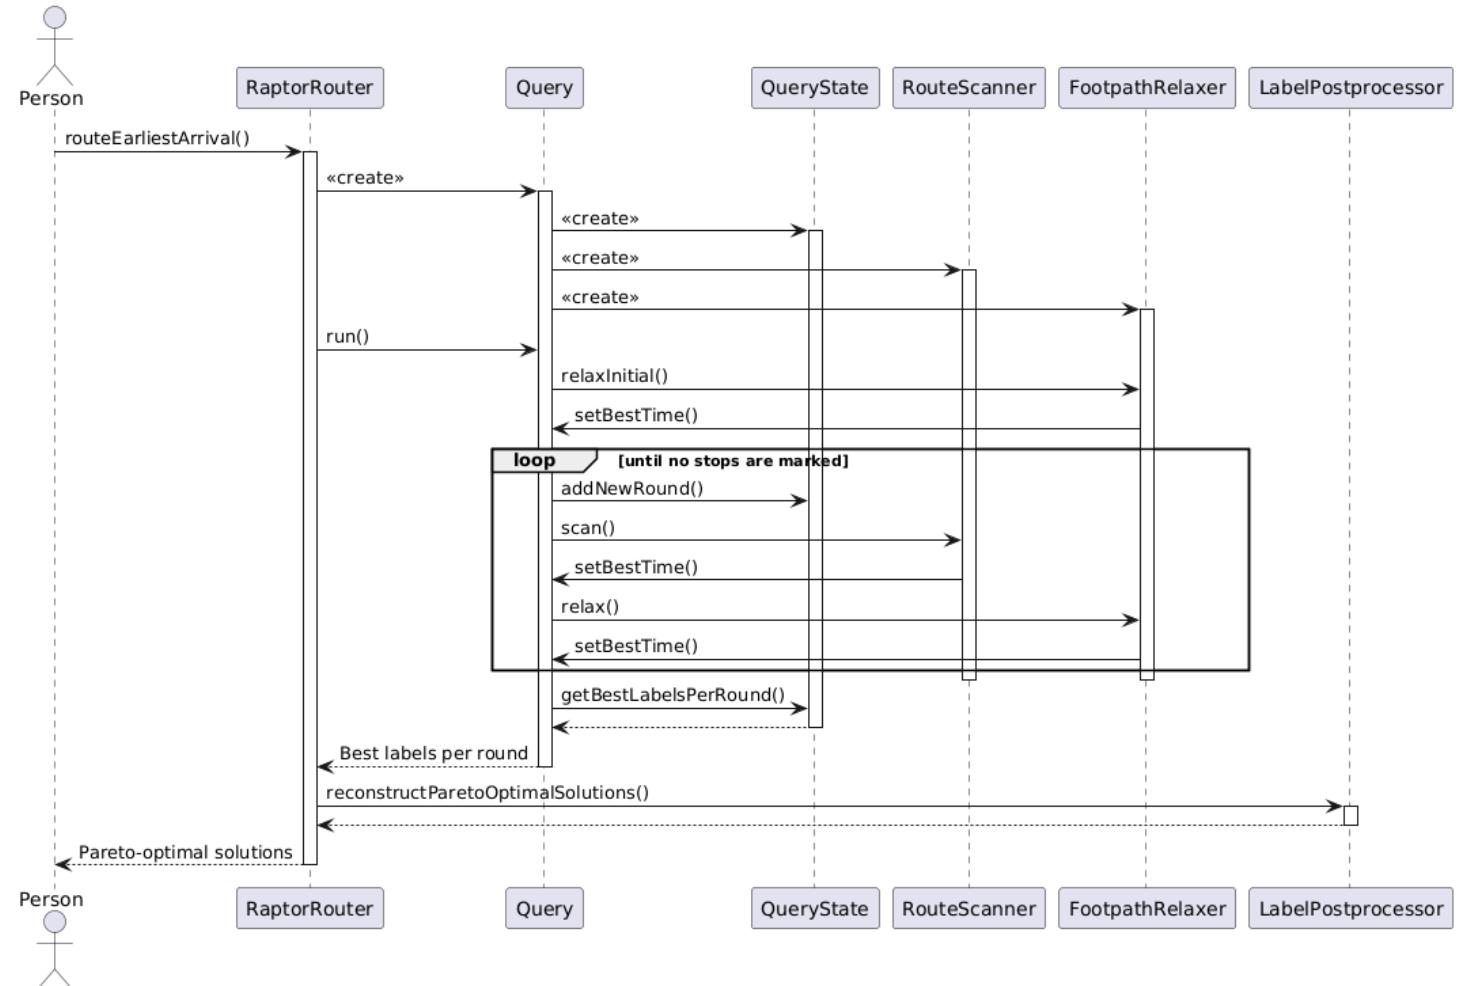
RAPTOR – Extended Features

- **Multi-stop routing:** Routes from multiple origins to multiple destinations, factoring in walking times
- **Reverse routing:** Optimize for latest departure instead of earliest arrival times
- **Multi-day routing:** Supports routing across multiple service days
- **Customizable queries:**
 - Transport mode, max transfers, walking limits
 - Wheelchair accessibility and bike accommodations
- **Range queries:** Find routes with shortest duration within specified departure windows
- **Isolines:** Calculate accessibility within a time range without predefined destinations
- **Performance:** Optimized stop time arrays for better CPU cache efficiency, no hashsets
- **Caching:** LRU eviction cache for masked stop times to enhance performance

Implementation

Extended RAPTOR – Query Route Calculation

- **Initialization:** Mark first stop
- **Initial footpath relaxation**
- **Main Loop: Round-based search**
 - Add new round
 - **Scan** routes, mark improved stops
 - **Relax** footpaths, mark improved stops
- **Termination:** End when no improvements possible
- **Reconstruct:** Best route



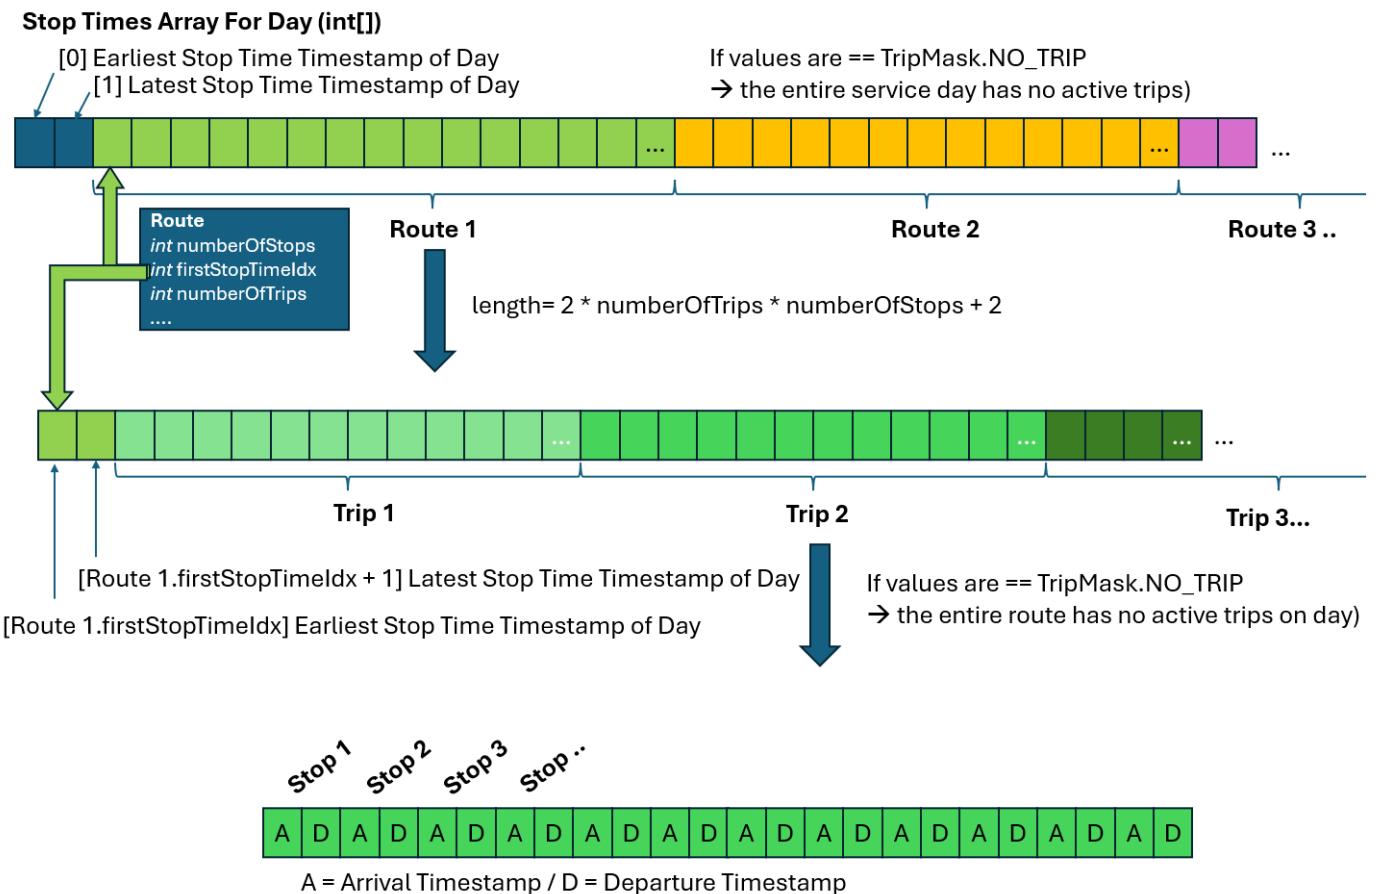
Implementation

Extended RAPTOR – Stop Times Array

- **Indexing** a stop time:

```
int index = route.firstStopTimeIdx  
+ 2 + 2 * (tripOffset * numberOfStops)  
+ stopOffset + typeConstant
```

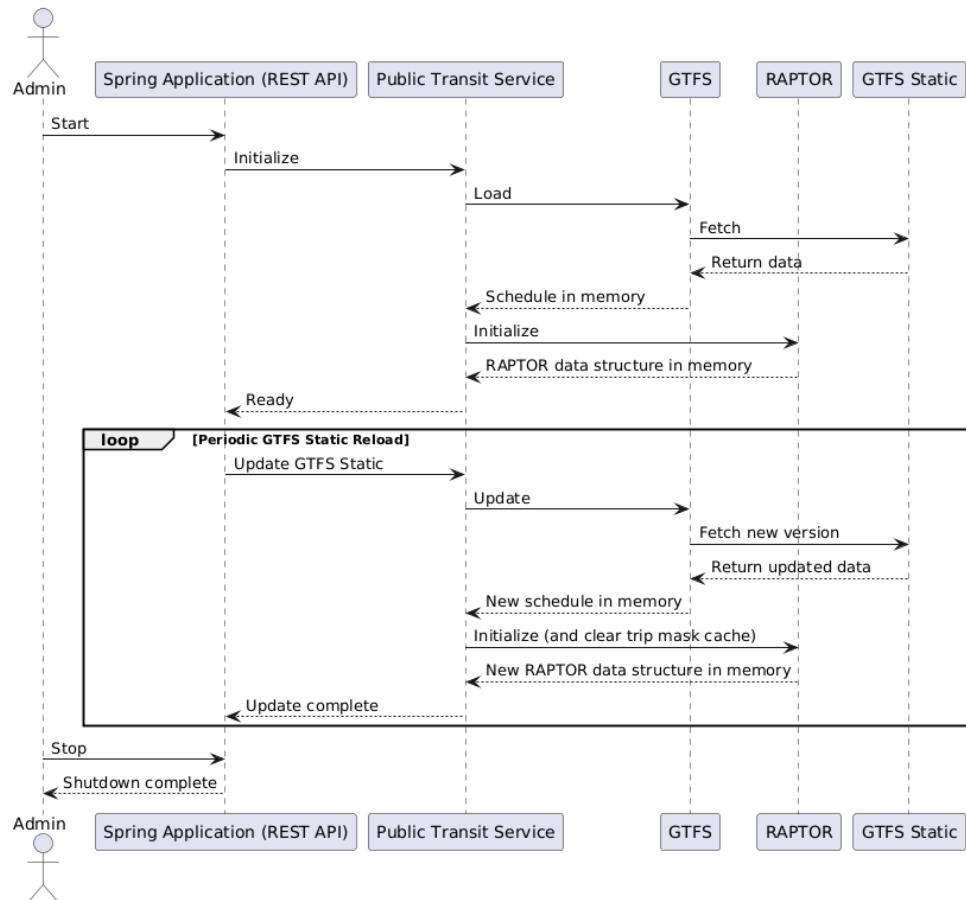
- tripOffset: Which trip number should be retrieved (can be incremented for the number of trips from route object)
 - number0fStops: Represents the number of stops each trip has (same for a given route)
 - stopOffset: Which stop number the time should be retrieved for
 - typeConstant: 0 if arrival time, 1 if departure time
 - +2: Skip earliest and latest trip time of the route
 - *2: Because every stop on the route has two values (arrival/departure)



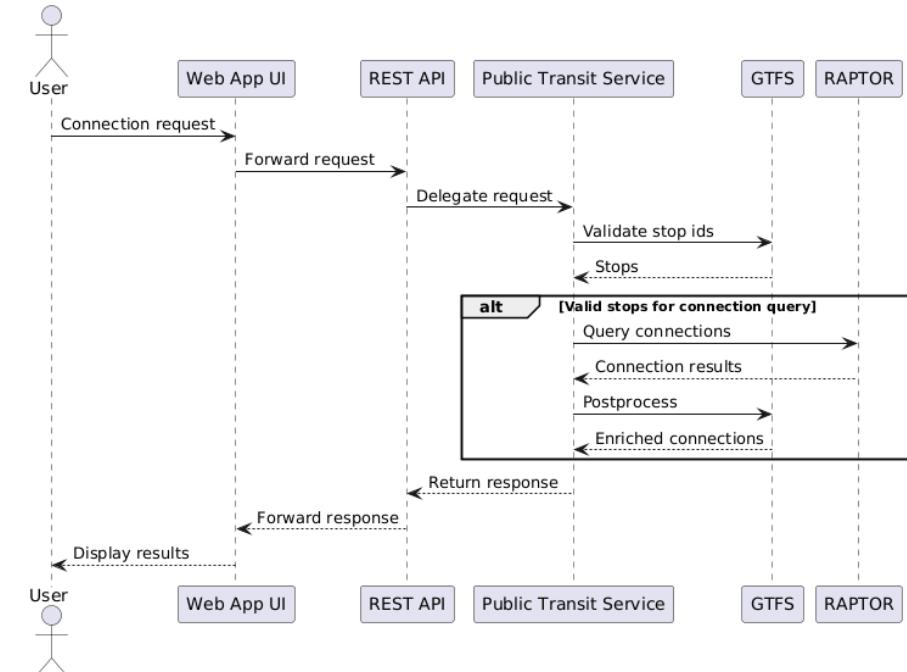
Implementation

Public Transit (Spring) Service

Application startup

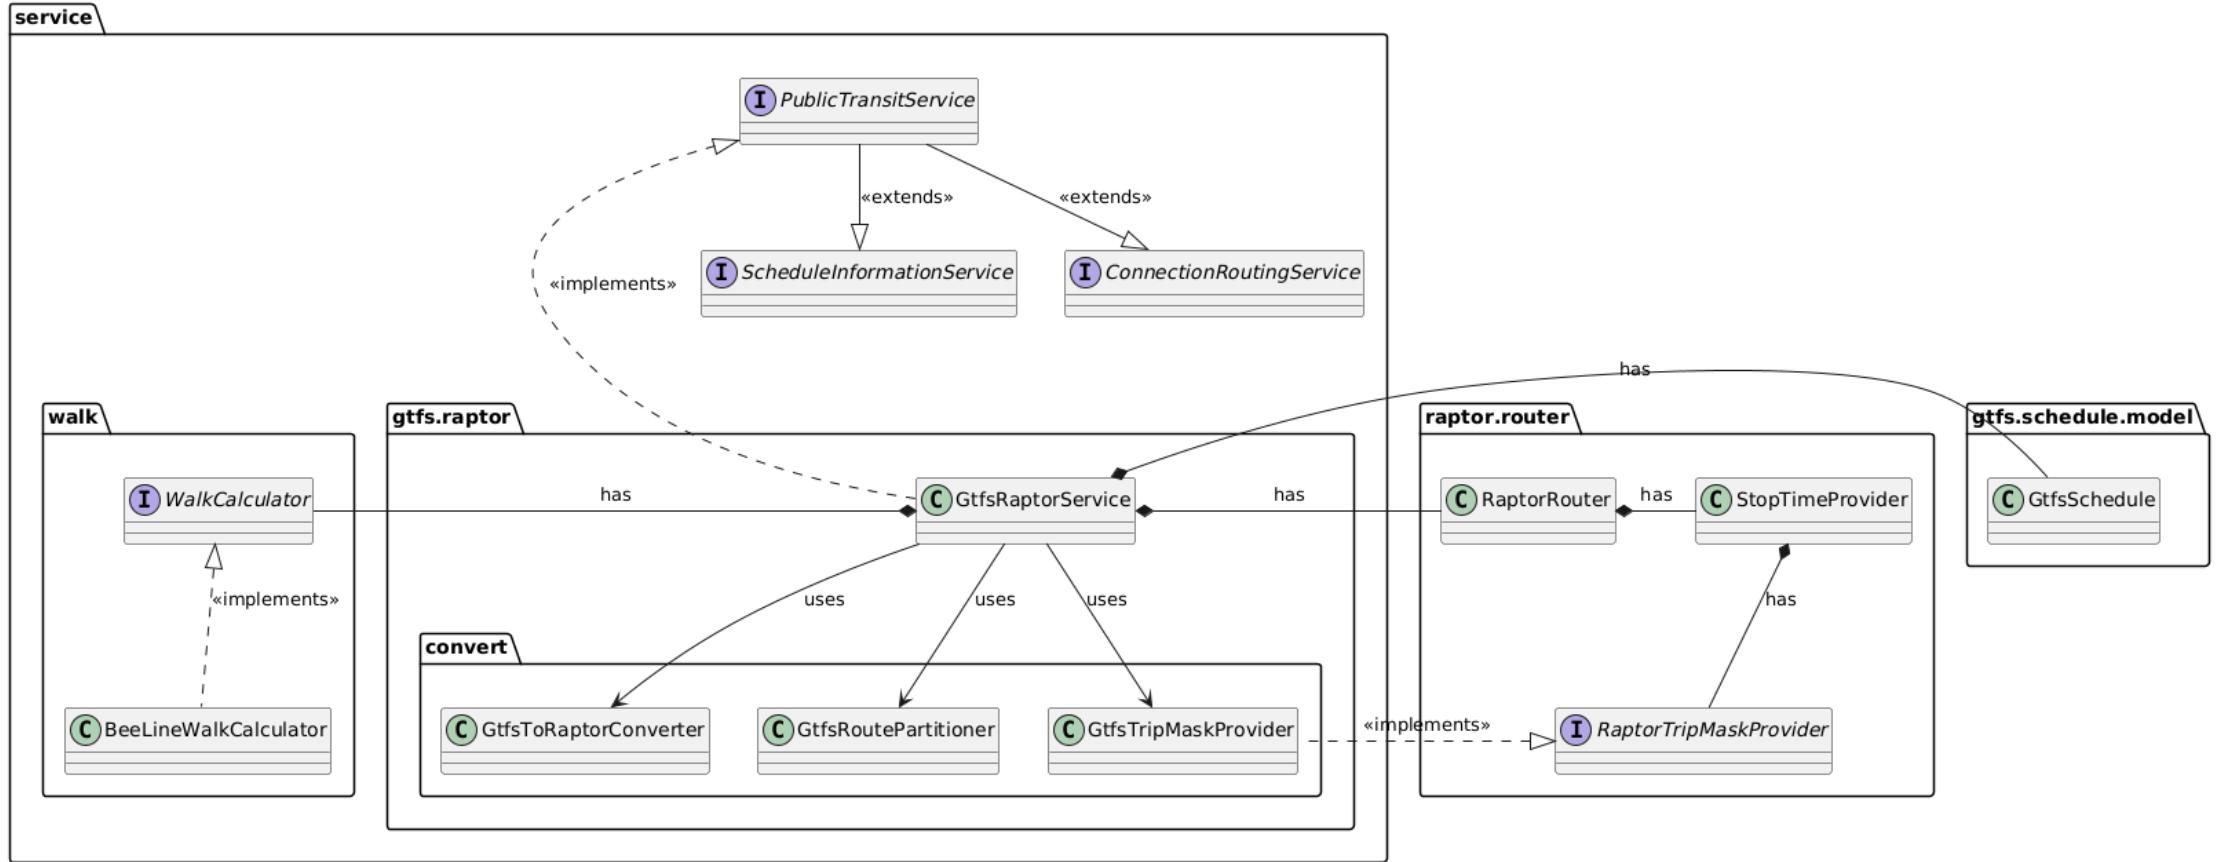


Connection request



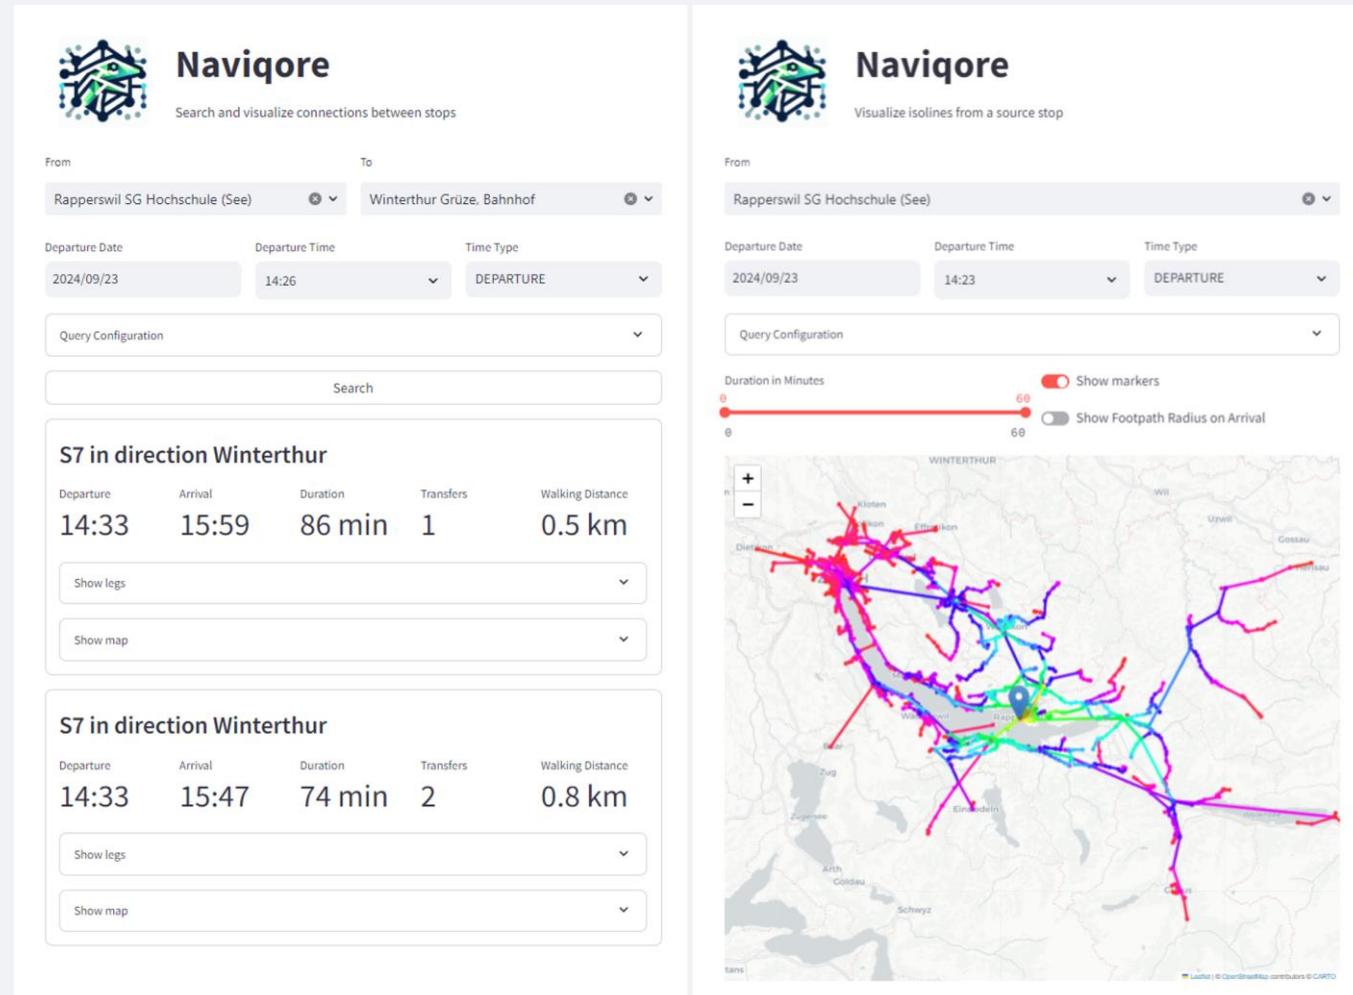
Implementation

Public Transit (Spring) Service



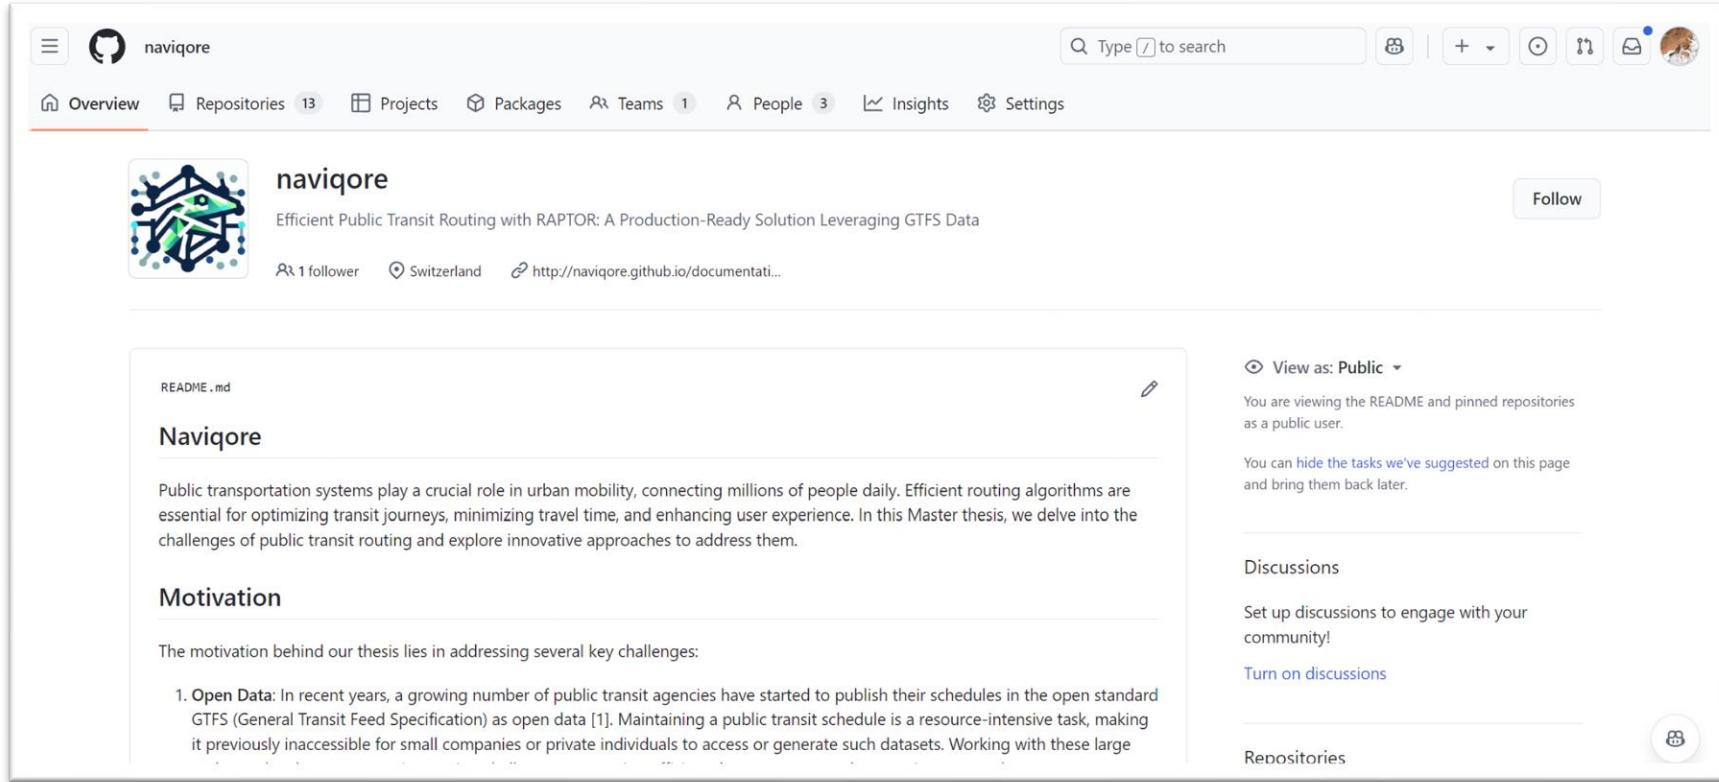
Results

Demo NAVIQORE



Results

Open-source Software



The screenshot shows the GitHub profile of the user 'naviqore'. The profile page includes a search bar, navigation links for Overview, Repositories (13), Projects, Packages, Teams (1), People (3), Insights, and Settings. The user has 1 follower and is located in Switzerland. A link to the documentation is provided.

naviqore
Efficient Public Transit Routing with RAPTOR: A Production-Ready Solution Leveraging GTFS Data
1 follower Switzerland http://naviqore.github.io/documentati...

README .md

Naviqore

Public transportation systems play a crucial role in urban mobility, connecting millions of people daily. Efficient routing algorithms are essential for optimizing transit journeys, minimizing travel time, and enhancing user experience. In this Master thesis, we delve into the challenges of public transit routing and explore innovative approaches to address them.

Motivation

The motivation behind our thesis lies in addressing several key challenges:

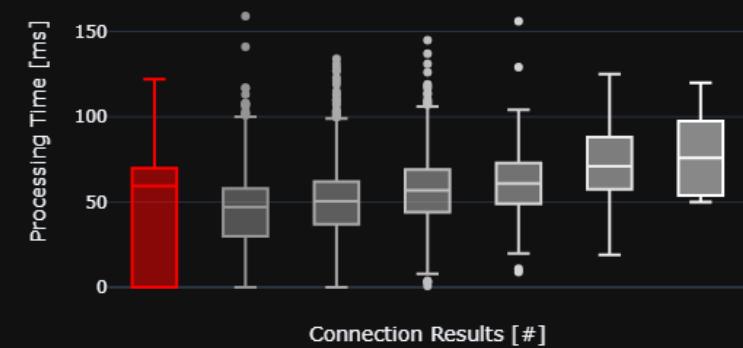
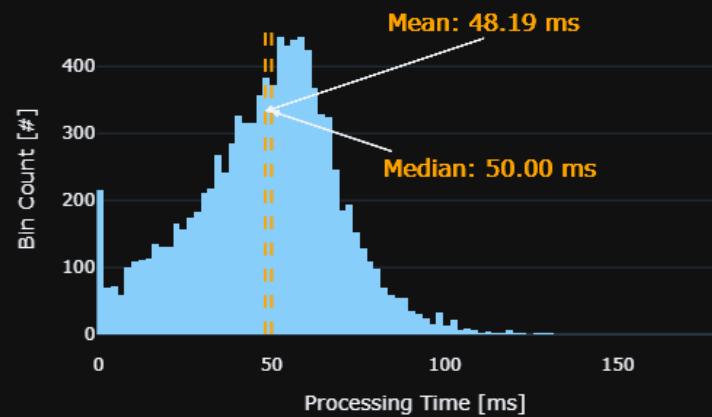
1. Open Data: In recent years, a growing number of public transit agencies have started to publish their schedules in the open standard GTFS (General Transit Feed Specification) as open data [1]. Maintaining a public transit schedule is a resource-intensive task, making it previously inaccessible for small companies or private individuals to access or generate such datasets. Working with these large

[naviqore/public-transit-service](#)

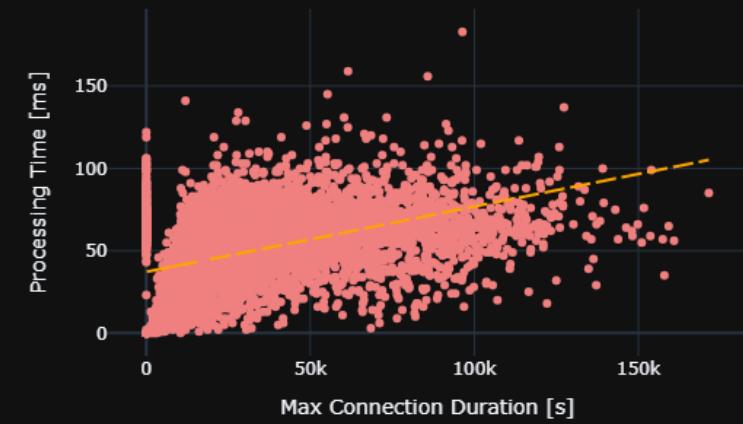
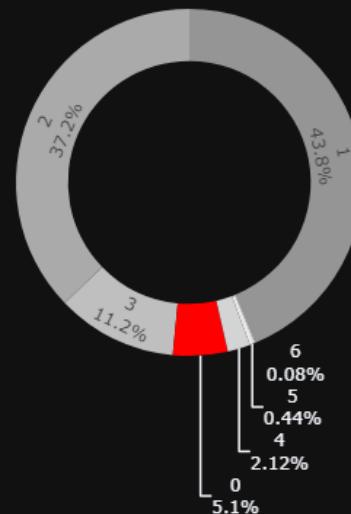
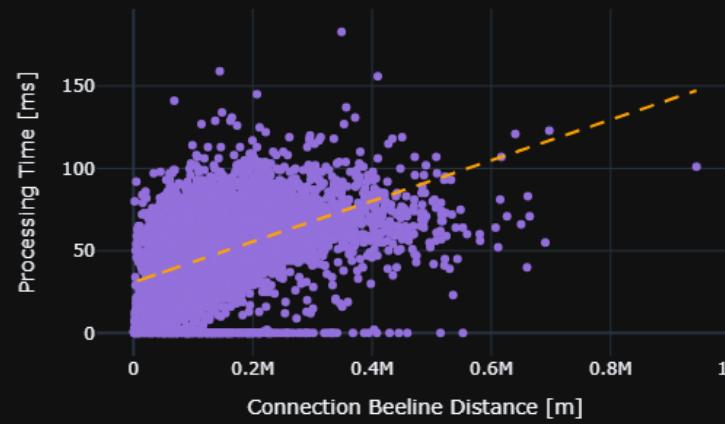
[naviqore/public-transit-client](#)

[naviqore/public-transit-viewer](#)

[naviqore/raptorxx](#)

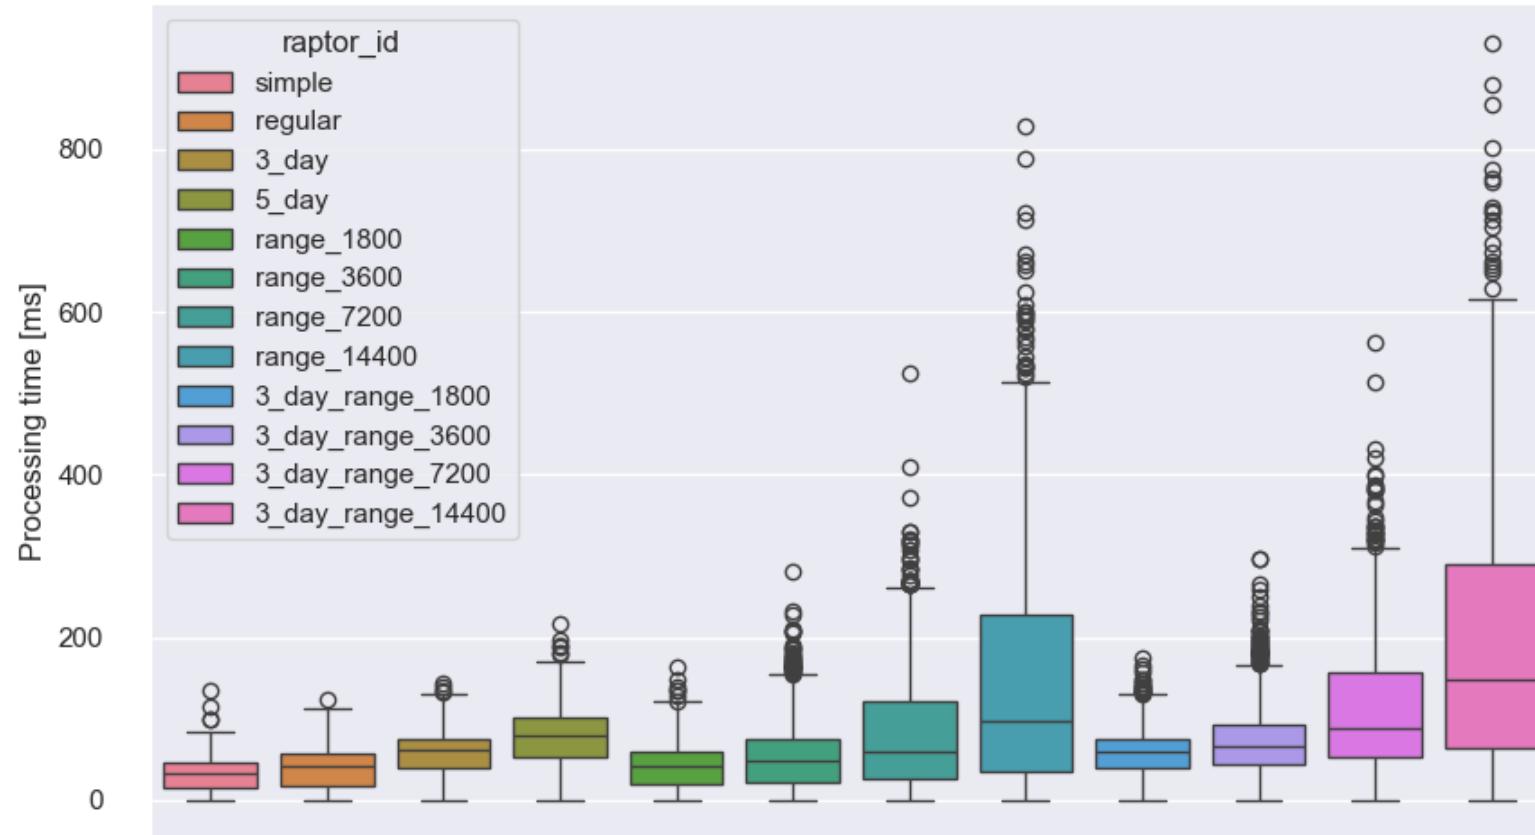


UC-SE-M3 - Efficiently request (in mean <250ms) connections between two stops.



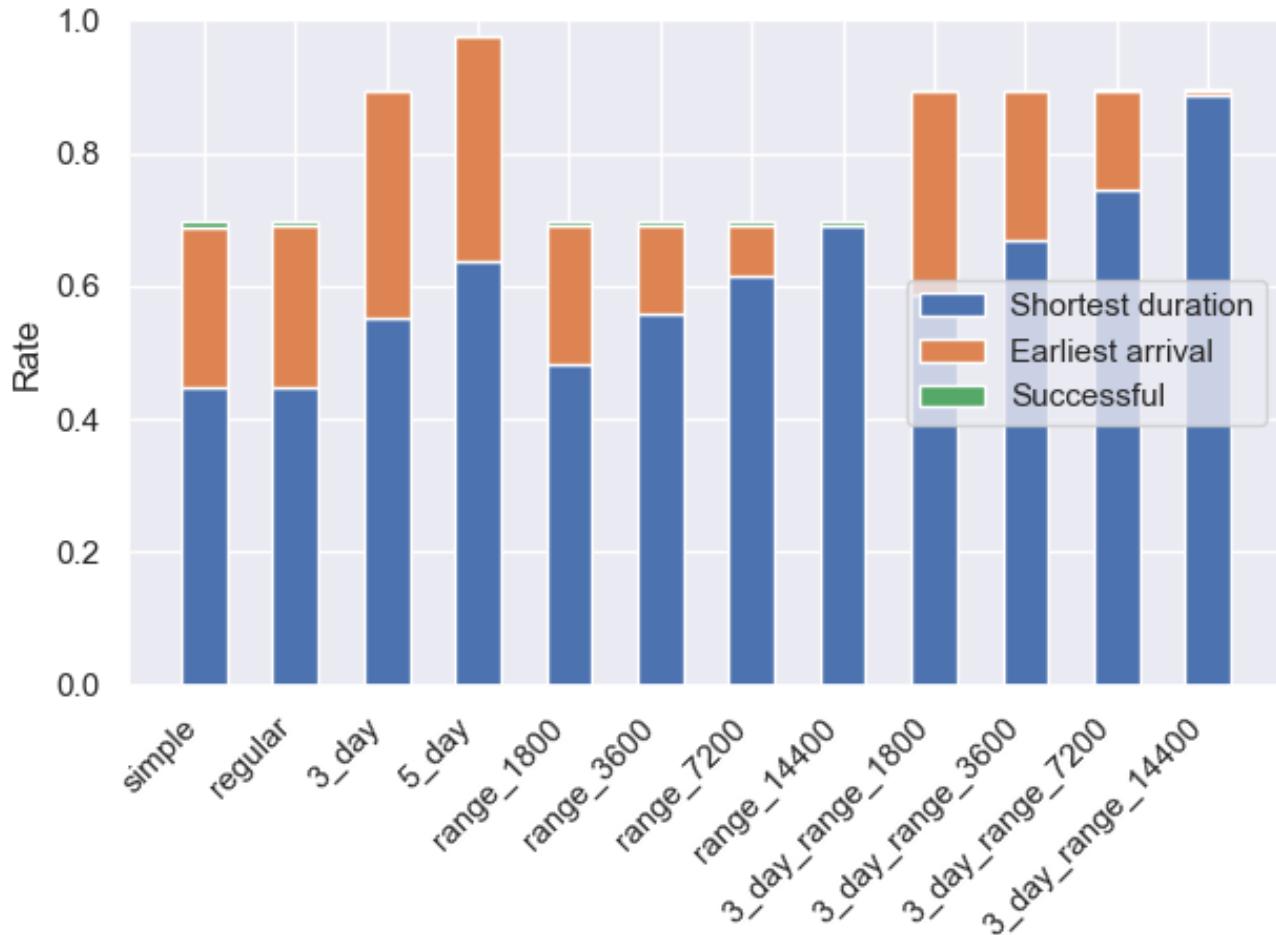
Results

Benchmark – Java RAPTOR: Performance



Results

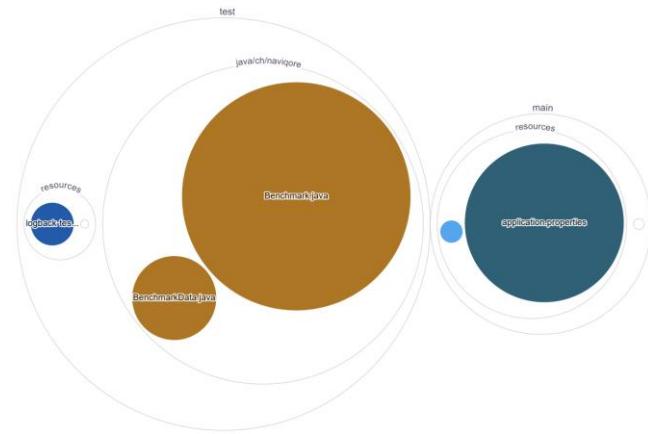
Benchmark – Java RAPTOR: Success Rate



Results

Code Metrics

- **Coding:** ~ 1300 commits, ~ 250 pull requests
- **Testing:** ~ 650 test cases, ~ 85% coverage

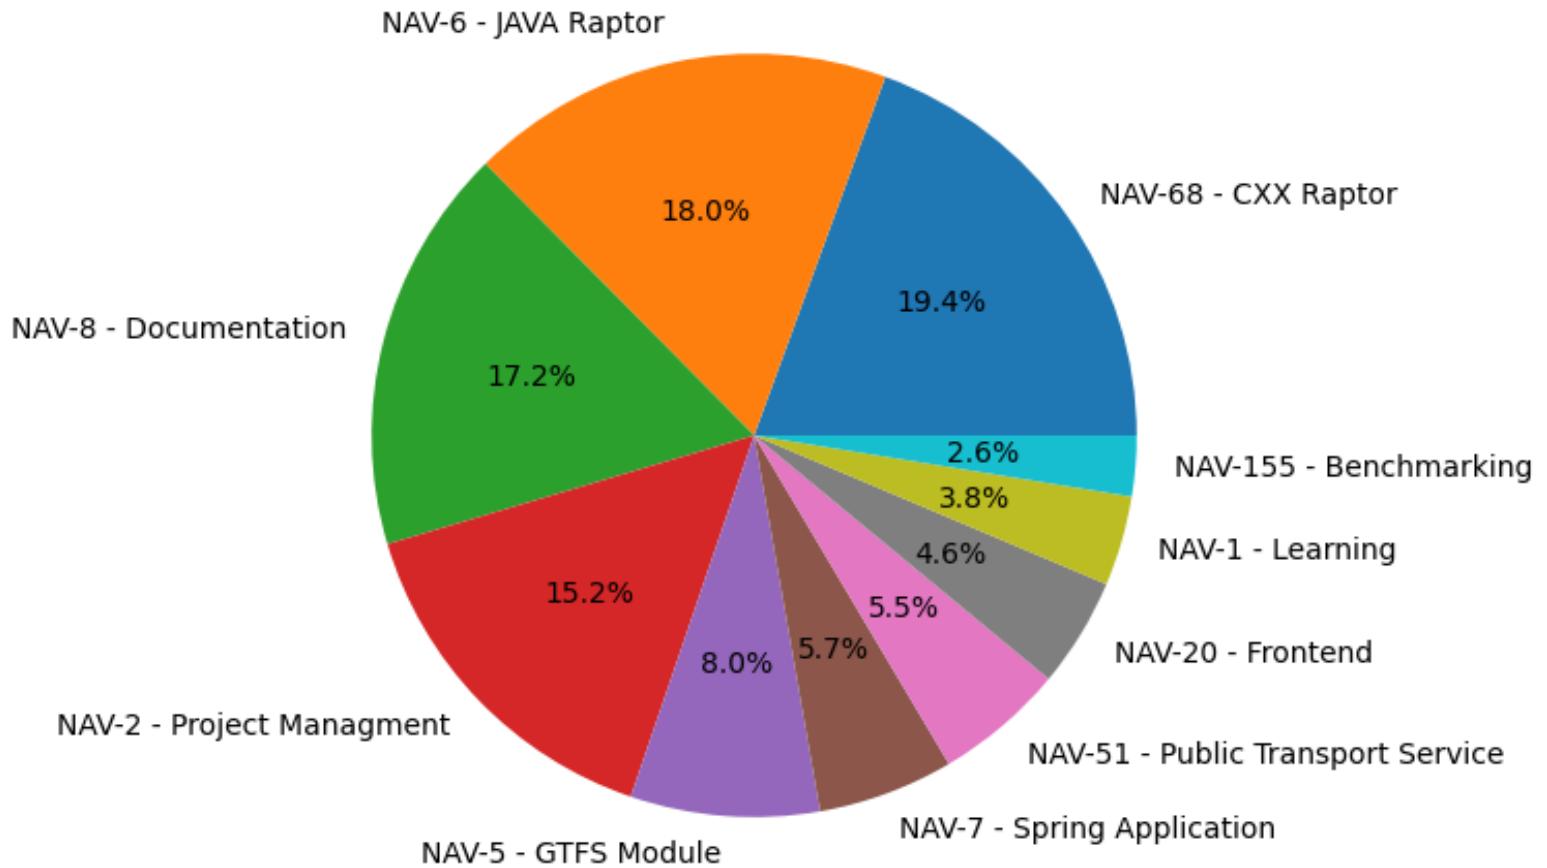


Repository Name	Language	Commits	Lines of Code (Test)	Test Cases (Integration)	Test Coverage
public-transit-service	Java	751	12116 (5662)	569 (51)	86%
public-transit-client	Python	72	822 (419)	23 (15)	80%
public-transit-viewer	Python	125	1007 (21)	1 (1)	-
raptorxx	C++	264	9164 (1018)	47 (0)	-

- **Notes:**
 - No coverage for **public-transit-viewer** due to Streamlit subprocess
 - No coverage for **raptorxx** due to MSVC and CMake limitations

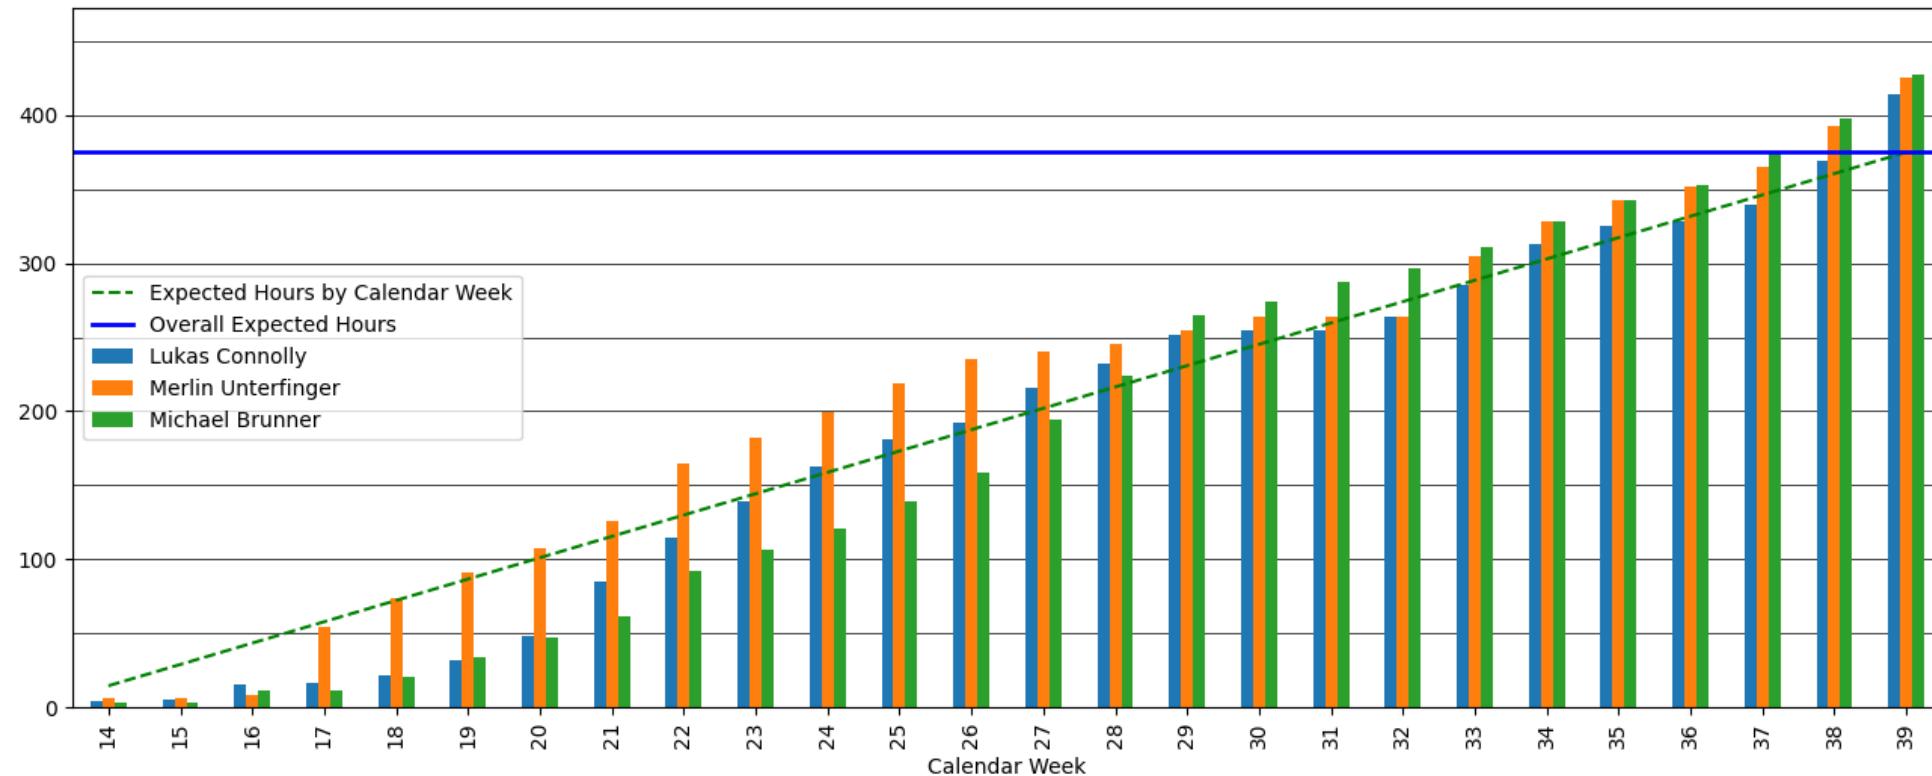
Results

Time Tracking – Epics



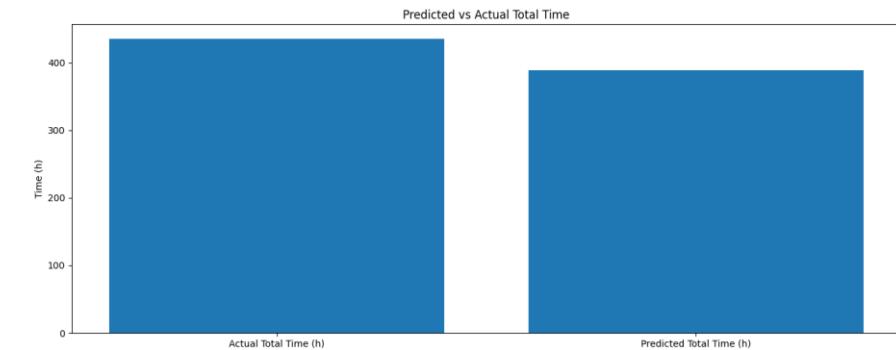
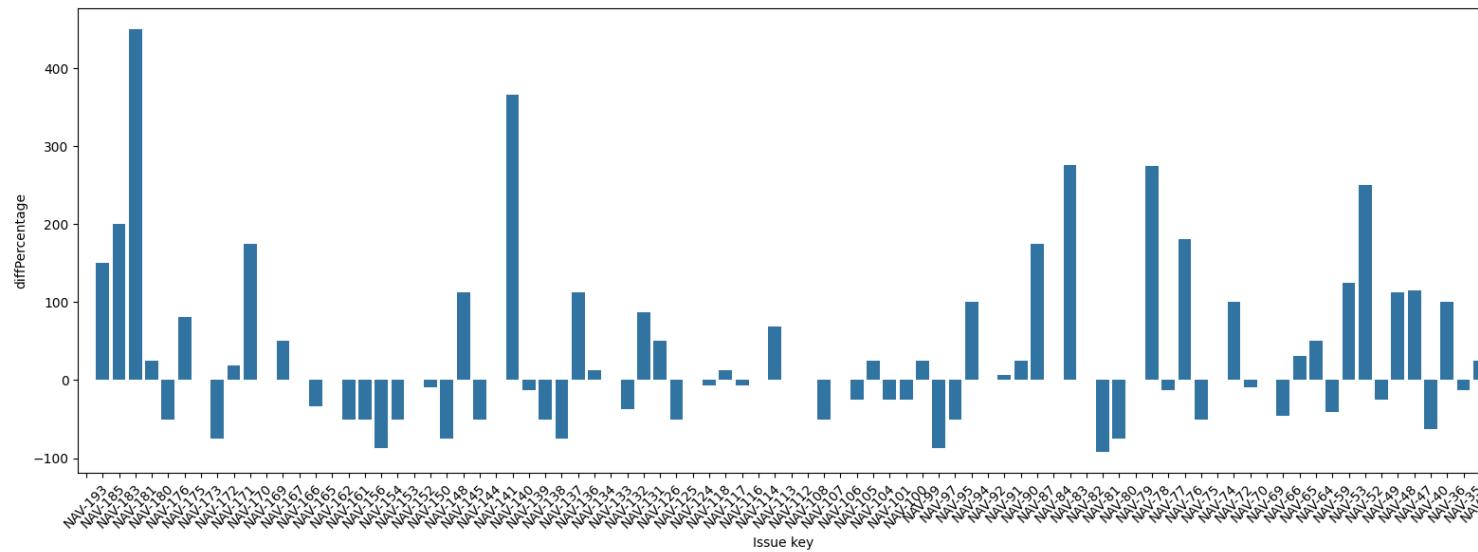
Results

Time Tracking – Developers



Results

Time Tracking – Estimates



Results

C++ vs. Java – Simple RAPTOR

- Motivation: Performance comparison
- Challenges:
 - CMake (build system)
 - Different OS / compiler
 - Memory management
- Results:**
 - Short distances slower
 - Long distances faster
- Our Java implementation uses primitive array, C++ dynamic std::vector

Source Stop	Target Stop	Iterations	Java Elapsed Time (ms)	C++ Elapsed Time (ms)	Average Difference (ms)	Average Difference (%)
8589640 (St. Gallen, Vonwil)	8579885 (Mels, Bahnhof)	100	13	19	-6	-31.6%
8574563 (Maienfeld, Bahnhof)	8587276 (Biel/Bienne, Taubenloch)	100	55	38	17	44.7%
8588524 (Sion, Hôpital Sud)	8508896 (Stans, Bahnhof)	100	45	20	25	125.0%
8510709 (Lugano, Via Domenico Fontana)	8579255 (Lausanne, Pont-de-Chailly)	100	46	29	17	58.6%
8574848 (Davos Dorf, Bahnhof)	8576079 (Rapperswil SG, Sonnenhof)	100	12	15	-3	-20.0%

Table 6.4: Comparison of Java and C++ performance.



Results

High-performance C++ Code → A Double-edged Sword

- Secure programming vs. performance?
- Raw pointer is **1.51 x faster** than shared pointer
- Raw array is
 - **3.09 x faster**
 - with `reserve(SIZE)` **1.41 x faster** than `std::vector` for adding/getting items



Benchmark	Time	CPU	Iterations
BM_raw_reference	367484 ns	325239 ns	2306
BM_shared_reference	495431 ns	452080 ns	1659
BM_vector	88914 ns	78125 ns	11200
BM_raw_array	125377 ns	112997 ns	4978

```
● ● ●

constexpr int SIZE = 10'000;

static void BM_raw_reference(benchmark::State &state)
{
    for (auto _ : state) {
        auto vec = std::vector<Song *>();
        vec.reserve(SIZE);

        std::ranges::transform(std::ranges::views::iota(0, SIZE),
                             std::back_inserter(vec),
                             []()<int> {
                                return new Song{"artist", "title"};
                            });

        std::ranges::for_each(vec, [](const auto *f) { delete f; });
    }
}

static void BM_shared_reference(benchmark::State &state)
{
    for (auto _ : state) {
        auto vec = std::vector<std::shared_ptr<Song>>();
        vec.reserve(SIZE);

        std::ranges::transform(std::ranges::views::iota(0, SIZE),
                             std::back_inserter(vec),
                             []()<int> {
                                return std::make_shared<Song>("artist",
                                                               "title");
                            });

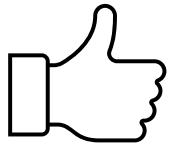
        vec.clear();
    }
}

BENCHMARK(BM_raw_reference);
BENCHMARK(BM_shared_reference);
```

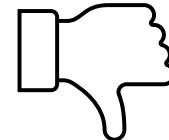
Limitations

- GTFS publications often lack information on **accessibility, bike allowance**, and pathways
- **International routing:** Combining multiple GTFS publications (duplicates?, time zones)
- Current **footpath routing:** Approximation using beeline distance and factor
- Trade-off: **Memory vs. performance** (caching)
- RAPTOR and **C++:** FFM API complexity, multiple OS, ...

Conclusion



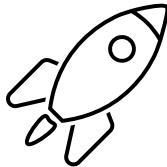
- Final product, **open source**
- Cool **team** and topic
- Software development process
 - Efficiency and **reviews**
- **Knowledge** gain
 - Data structures & algorithms, language specifics, architecture



- Little overlap of C++ and Java part of the project
- Work-life balance:
Strain on personal life...

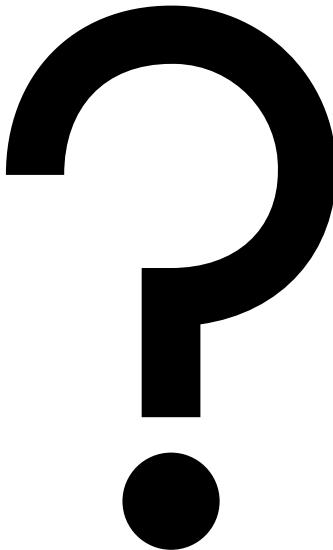
Results

Outlook



- Support **multi-criteria queries**, *mcRAPTOR*: Fares, journey attractivity, ...
- Incorporate **GTFS Realtime**
 - Consume protocol buffers
 - Update GTFS schedule and invalidate cache
- Implement **A*** algorithm for **routing footpaths**
 - Use OSM data to build the graph
 - Additionally: Calculate shapes of transit legs
- **Decoupling** of service and schedule / RAPTOR data
 - Common storage layer: Database holding GTFS and potentially RAPTOR data
 - Improves horizontal scaling

Any Questions



Results

References

- [1] “Gtfs cookbook,” <https://opentransportdata.swiss/de/cookbook/gtfs>, n.d., retrieved May 25, 2024.
- [2] M. Unterfinger, M. Brunner, and L. Connolly, “Efficient public transit routing with raptor: A production-ready solution leveraging gtfs data,” <https://github.com/naviqore>, 2024, retrieved September 26, 2024.
- [3] “General transit feed specification,” <https://gtfs.org>, n.d., retrieved May 25, 2024.
- [4] D. Delling, T. Pajor, and R. F. Werneck, “Round-based public transit routing,” in *Proceedings of the Meeting on Algorithm Engineering and Experiments (ALENEX)*. SIAM, 2012, pp. 130–140. [Online]. Available: <https://doi.org/10.1137/1.9781611972924.13>
- [5] K. Schwaber and J. Sutherland, “The scrum guide: The definitive guide to scrum: The rules of the game,” <https://scrumguides.org/docs/scrumguide/v2020/2020-Scrum-Guide-US.pdf>, 2020, retrieved May 25, 2024.
- [6] J. Sauer, “Public transit journey planning,” Lecture slides, Karlsruhe Institute of Technology (KIT), 2021, retrieved May 21, 2024.
- [7] R. C. Martin, *Clean Architecture: A Craftsman’s Guide to Software Structure and Design*, 1st ed. USA: Prentice Hall Press, 2017. [Online]. Available: <https://doi.org/10.5555/3175742>
- [8] J. Palermo, “The onion architecture, part 1,” <https://jeffreypalermo.com/2008/07/the-onion-architecture-part-1>, 2008, retrieved June 8, 2024.

Results

References

- [9] “Spring boot reference documentation,” <https://spring.io/projects/spring-boot>, n.d., retrieved May 31, 2024.
- [10] S. Brown, “The c4 model for visualising software architecture,” <https://c4model.com>, n.d., retrieved June 14, 2024.
- [11] ——, “The c4 model for software architecture,” <https://www.infoq.com/articles/C4-architecture-model>, 2018, retrieved June 14, 2024.
- [12] “Haversine formula,” https://en.wikipedia.org/wiki/Haversine_formula, n.d., retrieved May 25, 2024.
- [13] M. T. Goodrich and R. Tamassia, *Algorithm Design and Applications*, 1st ed. Wiley, 2014. [Online]. Available: <https://www.wiley.com/en-us/Algorithm+Design+and+Applications-p-9781118335918>
- [14] A. Horni, K. Nagel, and K. W. Axhausen, *The Multi-Agent Transport Simulation MATSim*. London: Ubiquity Press, 2016. [Online]. Available: <https://doi.org/10.5334/baw>
- [15] M. Rieser, D. Métrailler, and J. Lieberherr, “Adding realism and efficiency to public transportation in MATSim,” in *Proceedings of the 18th Swiss Transport Research Conference (STRC)*, Monte Verità, Ascona, May 2018. [Online]. Available: https://www.strc.ch/2018/Metrailler_Lieberherr.pdf
- [16] OpenJDK, “Jep 454: Foreign function & memory api (second preview),” <https://openjdk.org/jeps/454>, 2023, retrieved September 17, 2024.
- [17] “Project lombok: Cleaner, faster java,” <https://projectlombok.org>, n.d., retrieved May 25, 2024.